

::ISUG

techcast series

# Archive Database Access in ASE

October 16, 2007

SYBASE®

::ISUG

techcast series

## Archive Database Access in ASE



**Mike Harrold,  
Executive Director  
International Sybase Users Group**

SYBASE®

::ISUG

techcast series

## Archive Database Access in ASE



**Graham Ivey,**  
**Senior Staff Engineer**  
**Sybase, Inc.**

SYBASE®

## Motivation – Offline DBCC

- **The ability to run DBCC checks directly on a database dump using some (offline) utility:**
  - Does not impact production
  - Is not constrained by a small maintenance window
  - Does not require additional disk space

## Motivation – Object Level Restore


- **The ability to recover a single object (like a table) from a database dump**
  - An unqualified delete or update is accidentally done on a table
  - Data corruption has occurred within a table (never!!)
- **Reloading an entire database dump impacts production because it is time consuming**
  - Particularly if only a small table needs to be restored

# Motivation – Load Sequence Verification

- **The ability to know that:**
  - The database dump can be loaded
  - The transaction logs can be successfully applied to the database
- **This will guarantee that the entire load sequence will work if there is a catastrophic data loss**

# General Solution

- **Archive Database Access (ADA)**
- **This provides the means of constructing (or virtualising) a database on top of a database dump**
- **The virtualised database is called an archive database**



## General Solution (cont)

- An archive database looks and behaves like a traditional *read-only* database
- The *majority* of the database pages are stored within a database dump that has been taken earlier with **DUMP DATABASE**
- **Similar to a standby database except**
  - Made available *much* quicker (no data copying by **LOAD DATABASE**)
  - Consumes little traditional disk storage



# Solution for Offline DBCC

- **An archive database looks like any read-only database**
- **DBCC checks may be run against it**
  - The database can be made available on any ASE that has access to the dump file and thus need not affect production in any way
  - The maintenance window can be as big as is needed
- **Caveat**
  - DBCC checks against an archive database are not truly offline
    - **There is no small, standalone utility that validates the contents of the dump**
    - **The utility is ASE itself!**

# Solution for Object Level Restore

- **Use SELECT INTO from the archive database into the target database:**  

```
select * into <production_db>..<table>  
from <archive_db>..<table>
```
- **Use CIS and SELECT INTO or bcp if the archive database is available on a server other than the target server**
- **Caveat**
  - This is not equivalent to a LOAD TABLE command
  - Index entries are not restored automatically. They are restored by virtue of inserting restored rows into the target table


# **Solution to Load Sequence Verification**

- **Like a standby database, transaction logs can be applied to an archive database**
- **Load sequence verification can therefore be done**
- **The ability to apply transaction logs allows for both DBCC checks and object-level restore to be done using an up-to-date version of the database**

# Components of an Archive Database

- **An archive database consists of:**
  - A database dump (the archive)
  - A traditional database (known as a *scratch database*) that hosts a new system table called `sysaltusages`
  - Rows in the `sysusages`, `sysdevices` and `sysattributes` tables in the master database
  - Some traditional disk storage used to store any modified or newly-allocated pages. This is called the *modified pages section*
  - Optionally, a transaction log dump

# Components of an Archive Database



# The Database Dump

- **The database dump:**
  - Only stores a subset of pages of the dumped database
  - Is used as a repository for unmodified pages
  - Read-only. Changes cannot be made to the database dump
- **ASE sees the database dump and its stripes as database devices (in sysdevices) that are useable by the archive database only**

# The Modified Pages Section

- **An archive database is read-only**
- **However, a certain amount of write activity is needed for permissible operations**
  - Database recovery in order to make the archive database consistent
  - DBCC commands that perform fixes so that coherent versions of tables can be restored
- **Modified and newly-allocated pages must be stored somewhere other than the database dump**
- **Every archive database is allocated an amount of traditional database storage for this. This is called the *modified pages section***

## **The Modified Pages Section (cont)**

- **The modified pages section is reflected as rows in the sysusages table in much the same way that space allocated to a traditional database is**
- **The modified pages section:**
  - Is allocated when the archive database is created
  - Can be extended when necessary, to accommodate more modified and newly-allocated pages
- **Together, the database dump and the modified pages section provide the storage for the archive database**



# Sysaltusages and the Scratch Database

- **sysal tusages is a new data-only locked system table (datarows)**
  - Like sysusages for a traditional database, it is used to perform logical-to-virtual page mapping for the archive database
  - Because of the highly fragmented nature of an archive database, it has many more rows than sysusages
  - As a result, it is not stored in the master database
  - Instead it is stored in a *scratch database*
- **The scratch database is a traditional database that simply hosts sysal tusages**

## Sysaltusages and the Scratch Database (cont)

- A database can be designated as a scratch database using:  
`sp_dbopti on <dbname>, "scratch database", "true"`
- It is preferable that the scratch database is dedicated to the purpose of hosting sysal tusages:
  - sysal tusages is quite volatile
  - Truncate log on checkpoint can be set on the scratch database
  - The scratch database can be dropped and re-created if needed
- **Threshold procedures can be used to manage space within the scratch database**

## **Sysaltusages and the Scratch Database (cont)**


- **Each archive database can only be assigned a single scratch database**
- **Multiple archive databases can use the same scratch database**
- **If you have a large number of archive databases, you may define multiple scratch databases**

# Sysaltusages Schema

Col umn_name	Type	Length	Null s
dbi d	small i nt	2	0
l ocati on	i nt	4	0
l start	i nt	4	0
si ze	i nt	4	0
vstart	i nt	4	0
vdevno	i nt	4	0
segmap	i nt	4	0
i ndex_name	i ndex_descri pti on		
csysal tusages	cl ustered, uni que l ocated on system		
i ndex keys			
-----			
dbi d, l ocati on, l start			

# Sysusages and Sysaltusages



Traditional Database



master.dbo.  
sysusages





Archive Database




scratchdb.dbo.  
sysaltusages

Traditional  
Disk Storage



Database  
Dump



# Creating an Archive Database

- **Create and designate a scratch database:**

```
create database scratchdb on default=20  
sp_dboption "scratchdb",  
    "scratch database", "true"
```

- **Create the archive database with  
the CREATE ARCHIVE DATABASE command**

```
create archive database <dbname>  
    on <device> = <size> [, <device> = <size>]  
    with scratch_database = <scratch_dbname>
```

- **Example, an archive database with a 100 MB modified  
pages section**

```
create archive database archivedb  
    on datadev3 = 100  
    with scratch_database = scratchdb
```

# Populating an Archive Database

- An archive database is populated using **LOAD DATABASE**
- The syntax is similar to that of loading a traditional **LOAD DATABASE**:

```
load database <dbname>  
    from <dump_device1>  
    [stripe on <dump_device2> ...]  
    [with norecovery]
```

- **For example:**

```
load database archivedb from "/dumps/060415.dmp"
```

## Populating an Archive Database (cont)

- **The stripes of the database dump are inserted into sysdevices as if they are database devices**
  - They are marked with a special status to indicate that they are used by an archive database
  - The devices are given automatically generated logical device names SYSDEV\$\_\_<device number>



## Populating an Archive Database (cont)

- **The data and log pages are not actually loaded. Backup Server is not used. Instead, ASE:**
  - Scans the dump building up a map of physically contiguous fragments of pages and their offsets within the dump
  - Creates the relevant rows in the sysaltusages table that reflect the logical-to-virtual page mapping for that contiguous fragment
  - Because the stripe device is seen by ASE as a device within sysdevices, the virtual page number of the fragment is like any other virtual page number in sysusages (in 12.5.x, an 8 bit device number plus a 24 bit offset, in 15.0.x, a 32 bit device number plus a 32 bit offset)

# Loading an Archive Database (Without Recovery)

- **The WITH NORECOVERY option of the LOAD DATABASE command ensures that recovery is not run**
- ***Not* supported for traditional database loads**
- **Allows a database dump to be made visible:**
  - In the quickest possible time (no recovery)
  - Using the least amount of space within the modified pages section (since very few pages are modified)
- **The database is brought online automatically**
- **Because recovery has not been run, the database may be both logically and transactionally inconsistent**
- **Accessible only by a user with sa\_role**

## **Loading an Archive Database (With Recovery)**

- **More normally, WITH NORECOVERY is not specified**
- **Recovery (redo pass) is run at the end of LOAD DATABASE**
- **This will result in modified and newly-allocated pages**
- **When a page is first modified (or newly-allocated) it is automatically remapped to the modified pages section**

# Bringing the Archive Database Online

- **ONLINE DATABASE** brings the archive database online after the load  
`online database <dbname>`
- Recovery is run (undo pass) as is the case for (normal) **ONLINE DATABASE**)
- This will result in modified and newly-allocated pages
- When a page is first modified (or newly-allocated) it is automatically remapped to the modified pages section

# Loading a Transaction Log

- The syntax for LOAD TRAN is the same as that for a traditional database:  

```
load tran <dbname>  
      from <dump_device1>  
      [stripe on <dump_device2> ...]
```
- For example:  

```
load tran archivedb from "/dumps/060415_01.dmp"
```
- Like LOAD DATABASE, the stripes of a transaction log dump are inserted into sysdevices
- Unused devices from the previous LOAD TRAN are removed

## Loading a Transaction Log (cont)

- The recovery redo pass is run on the database. Any newly-allocated pages and pages modified for the first time are remapped to the modified pages section
- Unlike **LOAD TRAN** for a traditional database, **LOAD TRAN** for an archive database can follow an **ONLINE DATABASE** command
  - This is to allow the archive database to be useable throughout the load sequence for running DBCC checks and performing object-level restore

# Dropping the Archive Database

- **An archive database can be dropped using the DROP DATABASE command:**  
`drop database <dbname>`
- **Rows are removed from:**
  - master..sysusages
  - master..sysattributes
  - master..sysdevices
  - master..sysdatabases
  - <scratch database>..sysaltusages

# Management of the Modified Pages Section

- A page is remapped to the modified pages section when it is first modified or newly-allocated
- The modified pages section is logically divided into two areas (or segments):
  - The permanent changes segment
  - The disposable changes segment
- **In *general*:**
  - Any page modified by LOAD DATABASE or LOAD TRAN is mapped to the permanent changes segment
  - Any page modified outside of LOAD DATABASE or LOAD TRAN is mapped to the disposable changes segment



## Management of the Modified Pages Section (cont)


- **A maximum of three versions of a page can exist**
  - Page 100 is in the dump
  - Page 100 (the version above) is modified by LOAD DATABASE recovery and the new image is stored in the permanent changes segment of the modified pages section
  - Page 100 (the version above) is modified by ONLINE DATABASE recovery and the new image is stored in the disposable changes segment of the modified pages section
- **LOAD TRAN is able to follow ONLINE DATABASE by disposing of all the pages stored in the disposable changes segment**

Dump


Permanent  
Changes Segment

Disposable  
Changes Segment

Database



Dump



Permanent  
Changes Segment

Disposable  
Changes Segment

Database

## Management of the Modified Pages Section (cont)

- **Allocation of space to the permanent and disposable changes is maintained automatically**
  - Disposable changes segment steals 256 pages at a time from the permanent changes segment
- **Described by rows in sysusages**
- **Any change results in an update of sysusages at the end of the command**
- **Results in some logging in the master database**

## Management of the Modified Pages Section (cont)

- **At the start of every LOAD DATABASE and LOAD TRAN, all the space in the modified pages section is reallocated to the permanent changes segment**
- **Use the 15715 trace flag to override this**
  - Can reduce sysusages updates at the expense of wasting some space

## Management of the Modified Pages Section (cont)

- **Difficult to size the modified pages section**
  - Depends on the number of (unique) pages modified and newly-allocated by LOAD DATABASE/LOAD TRAN recovery
  - Depends on the number of (unique) pages modified and newly-allocated outside of LOAD DATABASE/LOAD TRAN, even if that page has already been modified by LOAD DATABASE/LOAD TRAN
- **Experience will help in sizing the modified pages section**
  - Start small and the modified pages section can be extended where necessary
  - Dropping and re-loading the database to reduce the size is not nearly as onerous a task as for a traditional database

## Extending the Modified Pages Section

- The **ALTER DATABASE** command can be used to increase the size of the modified pages section:  

```
alter database <dbname>  
on <device> = <size> [, <device> = <size>]
```
- The **ON** clause specifies the size and location of the expansion using traditional database storage
- Can be executed either when the archive database is idle or when it is suspended as a result of running out of space while executing the current command
- In this respect it is much like the suspension that happens when the last chance threshold is hit

## **DBCC Commands**

- **DBCC checks can be run against an archive database in the same way that it is run against any other database**
- **DBCC commands prevent other users from accessing an archive database while they are executing**
  - If you attempt to access an archive database while a DBCC command is being executed, you receive a message saying that the database is in single user mode



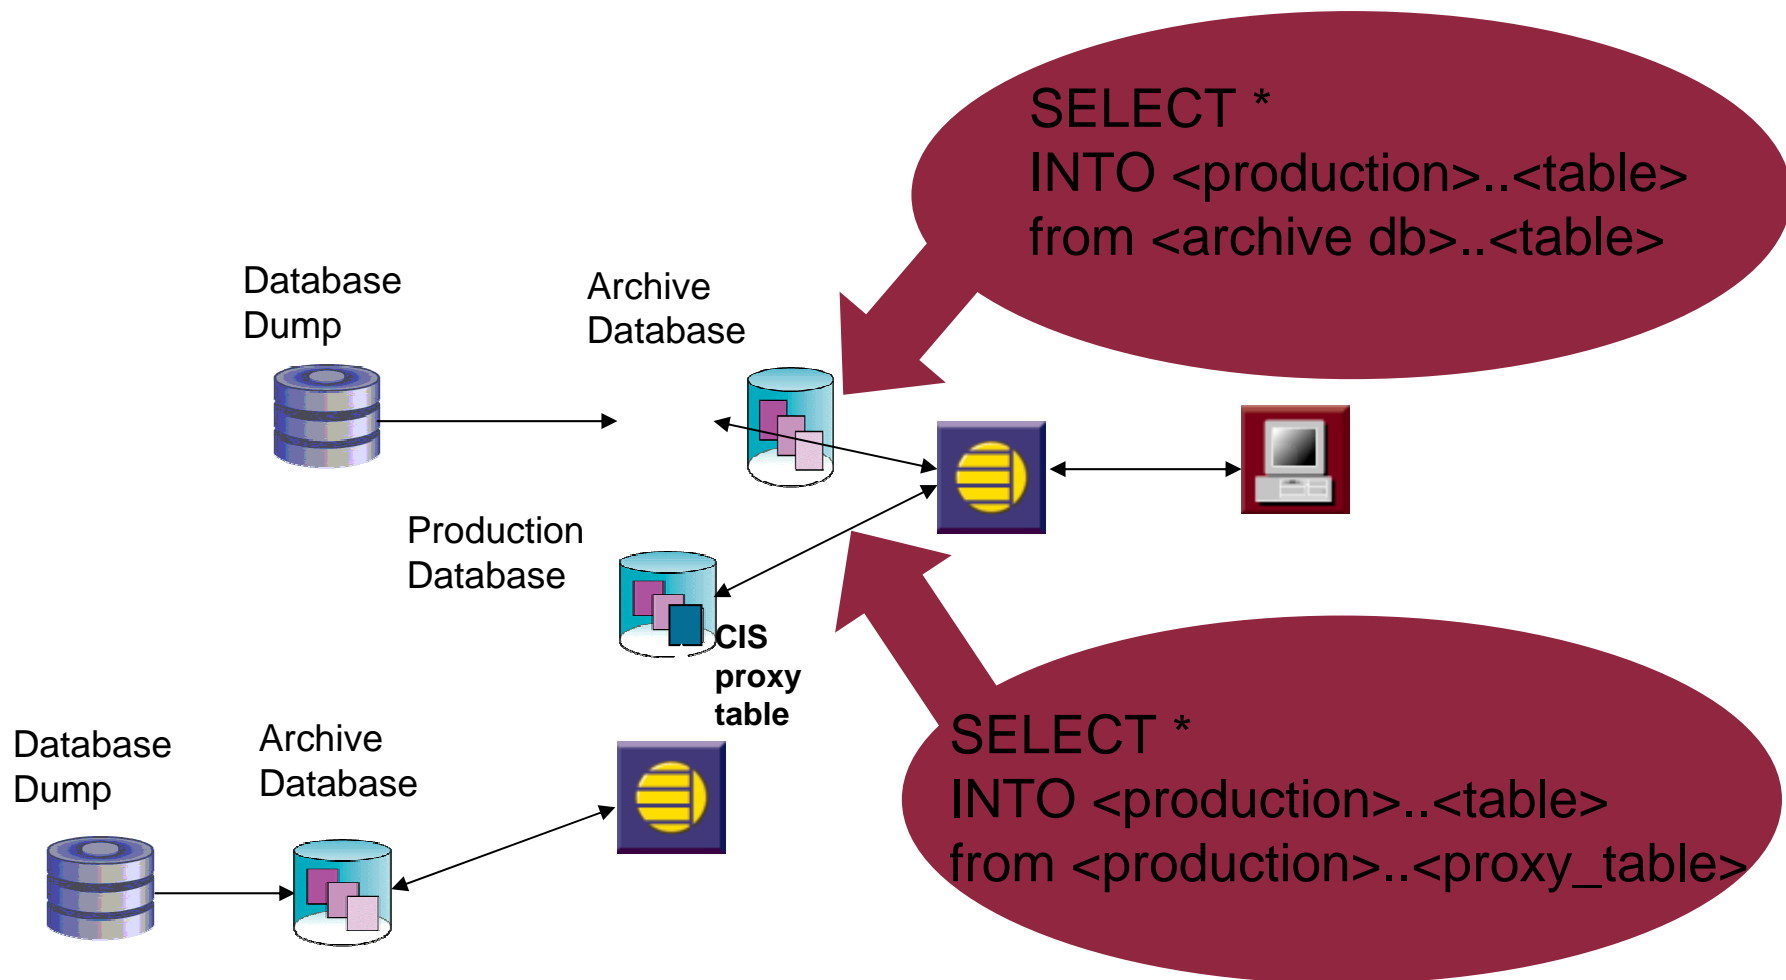
## **DBCC Commands (cont)**

- **The DBCC commands below are allowed in an archive database:**
  - checkstorage
  - checkdb
  - checkcatalog (the fix option is not supported)
  - checktable
  - checkindex
  - checkalloc
  - indexalloc
  - tablealloc
  - textalloc
- **The FIX option requires the database to be online**

## **DBCC Commands (cont)**

- **DBCC CHECKALLOC consumes a lot of space in the modified pages section**
  - Every allocation page (1 in 256) has information written to it
  - Modified pages are stored in the modified pages section
  - Therefore, the modified pages section must be at least 1/256th the size of the original database
  - True even if you run it with the NOFIX option

# Object Level Restore



# Compressed Dumps

- **Only dumps taken using:**  
DUMP DATABASE... WITH COMPRESSION  
DUMP TRAN ... WITH COMPRESSION  
**can be loaded into an archive database**
- **Dumps taken using the compress: : <n>: : <stri pe> syntax are not supported**
  - These dumps can be used by first decompressing the stripes using gunzi p
- **To load and use a compressed dump a special memory pool must be configured**

## Compressed Dumps (cont)

- When ASE reads a page from a compressed dump, it takes a compressed block from the dump, decompresses it, and extracts the required page. The decompression is done using large buffers from a special memory pool
- Each concurrent task uses a maximum of two buffers, one to read the compressed data and another to place the result of the decompression
  - By default, each buffer is 64 KB
  - Therefore, each concurrent task needs 64 (2K) pages
- **The size of the pool is configured using:**  
`sp_configure "compression memory size", <size>`  
<size> is given in 2 KB pages

## Compressed Dumps (cont)

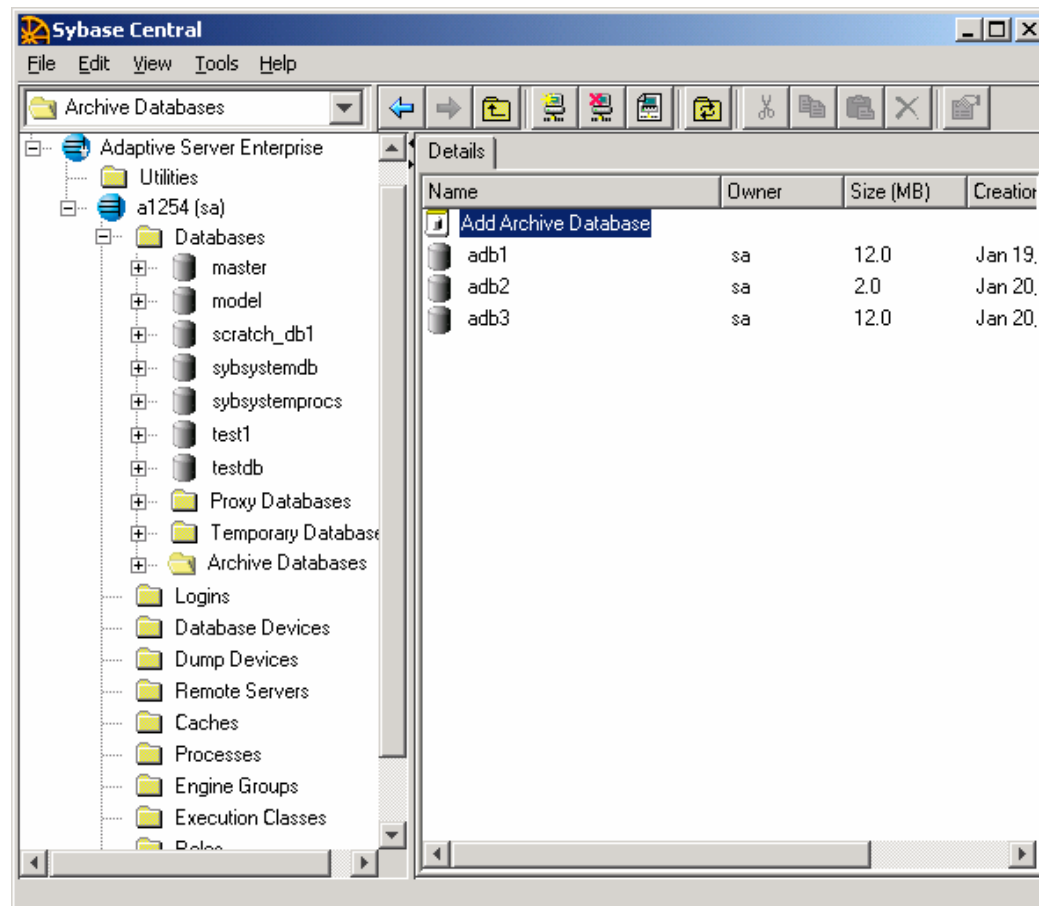
- **The compression format was changed to allow compressed dumps to be loaded**
- **The 12.5.4 Backup Server is the component that writes the new compression format. Therefore:**
  - Use a 12.5.4 Backup Server to load compressed dumps on a 12.5.3 (or earlier) installation if the dumps were created by a 12.5.4 Backup Server
  - Use a 12.5.4 Backup Server to create dumps on a 12.5.3 (or earlier) installation if the dumps are to be loaded into an archive database
- **The 12.5.4 Backup Server understands both 12.5.4 and earlier compression formats**

## Compressed Dumps (cont)

- **In 15.x, the 15.0.2 Backup Server is the component that writes the new compression format. Therefore:**
  - Use a 15.0.2 Backup Server to load compressed dumps on a 15.0.1 (or earlier) installation if the dumps were created by a 15.0.2 Backup Server
  - Use a 15.0.2 Backup Server to create dumps on a 15.0.1 (or earlier) installation if the dumps are to be loaded into an archive database
- **The 15.0.2 Backup Server understands both 15.0.2 and earlier compression formats**

# Sybase Central

- **Archive databases are supported in Sybase Central**





# Availability

- **12.5.4 and 15.0.2**
  - create archive database
  - load database (with and without recovery)
  - online database
  - alter database
  - drop database
  - select, etc
  - dbcc commands (except checkstorage and checkverify)
  - Compressed dumps
- **12.5.4 ESD 1 and 15.0.2**
  - dbcc checkstorage and checkverify
- **12.5.4 ESD 3 and 15.0.2**
  - load tran

- The “load” portion of LOAD DATABASE and LOAD TRAN is significantly faster for an archive database than a traditional database
- There is a small penalty (less than 15%) in accessing an uncompressed archive database
- The penalty for accessing a compressed archive database is *much* larger

## Performance (cont)

- **Case 1 (columns 2 and 3 in table)**
  - 7 GB database (~4 GB data and ~3 GB log)
  - Table being selected from had 11 million rows (700 MB)
- **Case 2 (columns 4 and 5 in table)**
  - 13 GB database (~11 GB data and ~2 GB log)
  - Table being selected from had 3 million rows (400 MB)

	Traditional	Archive	Traditional	Archive
load database	13 m	<<1 m	5 m	<<1 m
dbcc checkdb	34 m	33 m	211 m	179 m
dbcc checkalloc	16 m	13 m	50 m	49 m
dbcc checkstorage	11 m	20 m	43 m	40 m
select	5 m	5 m	88 s	99 s
select into	4 m	3 m	279 s	313 s
bcp (out)	14 s	13 s	174 s	199 s

# Archive Database Caveats and Limitations

- **Logging of sysusages updates in the master database**
- **Large amount of logging of sysal tusages updates in the scratch database**
  - Use a database exclusively for the scratch database if possible
  - Use truncate log on checkpoint
- **Increase the “number of devices” parameter to account for dump stripe devices**
- **Archive database is automatically put in single user mode when any command is run that changes the archive database – like DBCC commands**

## **Archive Database Caveats and Limitations (cont)**

- **Changes made by DBCC commands are not permanent**
  - Mapped to the disposable changes segment
  - Lost on next LOAD TRAN
- **Read-only**
- **No use of large buffer pools**
- **Default data cache only**
- **Maximum size of any stripe must be at most 32 GB (12.5.x only)**
- **No upgrade done. Across minor versions this is not generally an issue**

## **Archive Database Caveats and Limitations (cont)**

- **No free space thresholds**
- **Cannot be replicated**
- **No fail over in a High Availability system**
- **No DISK REFIT**
- **No cross-architecture load**
- **Tape dumps are not supported**
- **Remote dumps (using the AT syntax) are not supported**
- **Cannot specify the recovery order with `sp_dbrecovery_order`**

## Possible Future Additions

- **Allow an archive database that has been used to load an entire load sequence, to be *dumped***
  - No (or little) data copy involved in the dump phase
    - **Dump consists of the original database dump and the modified pages section**
  - The consolidated dump can be loaded into a traditional database
  - No need for a lengthy LOAD TRAN sequence
- **Allow UPDATE STATISTICS to be run**
  - Use *optdiag* to export statistics from archive database into production database