

::ISUG

techcast series

The ASE MDA Tables

Finding Out What Goes On Inside ASE

Rob Verschoor
Senior Evangelist
Sybase, Inc
robv@sybase.com

Peter Dorfman
Senior Staff Software Engineer
Sybase, Inc.
peter.dorfman@sybase.com

About Peter

- 12 years at Sybase ASE Engineering
- Designs and develops system management, monitoring and diagnostic tools for the ASE
- Leads the design and development of system management features for the ASE and other Sybase database products
- Peter.dorfman@sybase.com

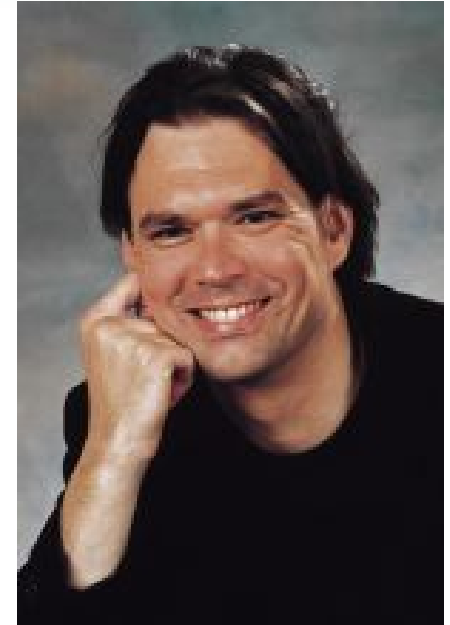


About Rob

- Consultant for ASE & Replication Server
- Worked with ASE since 1989
- Based in The Netherlands
- www.sypron.nl

Published three books about Sybase:

- “The Complete Sybase Replication Server Quick Reference Guide” (new)
- “The Complete Sybase ASE Quick Reference Guide” (3rd edition, new)
- “Tips, Tricks & Recipes for Sybase ASE”

**Since 1 February 2005: joined Sybase, Inc.**

- Sybase Engineering, Evangelism Group

- **Quick introduction to MDA tables**
- **Possible applications of MDA tables**
 - What's that application doing?
 - Identifying unused indexes
 - Identifying 'hot' tables
- **Historical MDA tables**
- **Archiving historical MDA table data**
- **Performance impact of MDA tables**
- **Counter wrap**
- **Analyzing stored procedure activity**
- **Miscellaneous topics**
- **Recent Enhancements**
- **Q&A**

Quick Introduction to MDA Tables

- **MDA tables were introduced in ASE 12.5.0.3**
 - MDA = Monitoring and Diagnostic Access
 - also called “monitoring tables”
- **35 proxy tables in ‘master’ database**
 - monSysSQLText, monObjectActivity, monCachedObject (etc.)
 - Can be accessed with regular SQL statements
 - When queried, tables are created on-the-fly from memory structures
 - No data is stored in **master** database
 - The proxy tables may also be created in a different database than **master**
- **Must be installed: ‘installmontables’ script**
- **No license needed: included in ASE base product**
- **Only small performance impact on ASE (<5%)**

Quick Introduction to MDA Tables

- **MDA tables provide access to low-level monitoring data**
 - Resource usage in ASE per table/query/entire server
 - Current activity in ASE per spid/query/procedure/table
 - Recent activity: recently completed statements, with the resources they required
- **Some examples of practically relevant information:**
 - Amount of memory occupied in the data cache by each table or index
 - Most frequently used tables/procedures
 - Top-N queries for CPU, I/O, elapsed time,...
 - Find unused indexes
 - SQL text of currently executing/recently executed statements
 - Automatically kill user processes that have been idle for more than X minutes
 - Provide server status information even when **tempdb** is full

Quick Introduction to MDA Tables

- **For more MDA basics, and a brief discussion of all tables:**
 - See Rob's presentation from Techwave 2003 (www.sypron.nl/mda)
- **In this presentation:**
 - We want to go one step further than just the basics
 - Look at practical applications of MDA tables
 - Things that are useful for you as a DBA

Possible Applications of MDA Tables

What's That Application Doing?

- **Does this sound familiar?**
 - A third-party 'black box' application runs on your ASE server
 - You have the feeling it sometime slows down the entire server...
 - ... but you don't know which queries it is sending to ASE
- **Classic solutions:**
 - Use "cmdtext" auditing to intercept the application's T-SQL commands
 - Use traceflag 11202 (writes all incoming client language to the errorlog)
 - Use third-party tools to find T-SQL commands by intercepting network packets
 - dbcc sqltext()
 - ...but all these methods have significant limitations or drawbacks

What's That Application Doing?

- **Solution: MDA tables monProcessSQLText & monSysSQLText**
- **monProcessSQLText: currently executing SQL**
- **monSysSQLText: recently executed SQL, now completed**
 - Historical table
 - Lets you “look back” in time
 - By copying rows regularly into an ‘archive’ table, complete history can be preserved

What's That Application Doing? Statement Statistics

- **Also: monSysStatement: info about completed SQL statement**
 - Number of logical I/Os
 - Number of physical I/Os
 - Number of network packets sent/received
 - Number of milliseconds of 'waiting time' during statement execution
 - Exact starttime & endtime of execution
 - Not the SQL Text itself; for this, see **monSysSQLText**
 - Historical table
 - Lets you "look back" in time
 - By copying rows regularly into an 'archive' table, complete history can be preserved

```
select * into #ts from master..monSysStatement
```

```
select KPID, BatchID, LineNumber, LogicalReads, Elapsed = datediff(ms,
    StartTime, EndTime) from #ts where LogicalReads > 10000 order by 1,2,3
```

KPID	BatchID	LineNumber	LogicalReads	Elapsed
-----	-----	-----	-----	-----
450298249	10	1	17844	0
452133791	3	1	40401	0
452790238	8	1	53003	0
453837392	10	1	17970	0
456197724	1	1	20585	0
456853857	9	1	49588	2000
457443052	8	1	66647	4000
458622875	19	1	22402	0
458622875	20	1	11214	1000
459014430	8	1	53146	0
459080255	1	1	108416	7000
459278288	16	1	17832	2000
460785077	9	1	55619	3000

```
select * into #tsql from master..monSysSQLText
```

```
select SQLText from #tsql where KPID= 459080255
order by BatchID, SequenceInBatch
```

SQLText

```
select  admnr = lvd.id_lm_adres,
        lvd.id_logmiddel
from
        logmi.dbo.lm_voorraad lvd,
        ravar.dbo.adm_relatie adm
where
        lvd.cdsys_lm_adrestype = "A"
and    lvd.cdsys_lm_opslagstat = "O"
and    lvd.id_lm_adres        = adm.admnr
and    adm.dat_ingang          <= @vandaag
and    (adm.dat_einde          >= @vandaag
or      adm.dat_einde          = null)
```

Monitoring Table Activity

- **Monitoring activity on a specific table**

```
select SQLText from master..monSysSQLText  
where SQLText like '%MyTable%'
```

- **Also handy for RepServer DBAs:**

- Quick way to figure out exactly which SQL is executed against your replicate DB
- Especially handy when developing/debugging custom function strings

Monitoring Index Utilization

- **Have you ever wanted to see**
 - Which indexes are never used?
 - How frequently they are used?
 - How many inserts, deletes, updates, physical or logical I/O they incur?
- **monOpenObjectActivity table provides:**
 - Table usage count
 - Index usage count
 - Last used dates
 - Physical, logical I/O
 - Row-level insert/delete/update counts
 - Lock wait counts for tables and indexes
- **NOTE: Statistics are reset when server is booted or object descriptor is reused in memory.**

monOpenObjectActivity

jisql oban:5003

Go Server: oban:5003 Database: master

Input window

```
select "Database" = db_name(DBID), "Table" = object_name(ObjectID, DBID), IndID = IndexID, UsedCount, LastUsedDate, OptSelectCount, La
from monOpenObjectActivity
order by UsedCount
```

Output window using table

Database	Table	IndID	UsedCount	LastUsedDate	OptSelectCount	LastOptSelectDate
master	spt_monitor	0	7	2004-01-24 16:05:05.15	7	2004-01-24 16:05:05
pubs2	titleauthor	0	8	2004-01-25 09:33:15.063	8	2004-01-25 09:33:15
tempdb	#1_____00000...	0	10	2004-01-24 15:46:57.156	8	2004-01-24 15:46:57
master	monCachedObject	0	13	2004-01-24 16:40:38.163	13	2004-01-24 16:40:38
master	monOpenObjectActivity	0	13	2004-01-25 09:55:41.063	13	2004-01-25 09:55:41
master	monSysStatement	0	16	2004-01-24 17:26:48.163	3	2004-01-24 15:33:07
master	spt_values	1	22	2004-01-24 15:31:03.07	23	2004-01-24 15:31:03
master	monProcessSQLText	0	36	2004-01-24 16:12:40.153	11	2004-01-24 16:03:14
pubs2	authors	1	36	2004-01-25 09:31:41.663	36	2004-01-25 09:31:41
pubs2	salesdetail	3	40	2004-01-25 09:55:37.063	40	2004-01-25 09:55:37.063
pubs2	sales	0	41	2004-01-25 09:51:40.066	41	2004-01-25 09:51:40.066
pubs2	title	1	50	2004-01-25 09:53:47.066	50	2004-01-25 09:53:47.066

Status

Table and
Index Usage

- Counts
- Dates

```
select "Database" = db_name(DBID), "Table" = object_name(ObjectID, DBID),
      IndID = IndexID, UsedCount, LastUsedDate, OptSelectCount, LastOptSelectDate
from master..monOpenObjectActivity
order by UsedCount
```


monOpenObjectActivity

jisql oban:5003

Go Server: oban:5003 Database: master

Input window

```
select "Database" = db_name(DBID), "Table" = object_name(ObjectID, DBID), IndexID, RowsInserted, RowsDeleted, RowsUpdated, LockWaits
from monOpenObjectActivity
order by RowsInserted desc
```

Output window using table

Database	Table	IndexID	RowsInserted	RowsDeleted	RowsUpdated	LockWaits
testdb	t1	0	63	51	11	1
tempdb	#16	00000120016207904	0	13	0	0
tempdb	#11	00000120016207904	0	12	0	0
tempdb	#12	00000120016207904	0	12	0	0
tempdb	#14	00000120016207904	0	12	0	0
tempdb	#15	00000120016207904	0	12	0	0
master	syblicenseslog	0	3	0	0	0
pubs2	sales	0	0	0	0	0
pubs2	titles	0	0	0	0	0

Status

Per Table

- Inserts
- Deletes
- Updates
- Lock Waits

```
select "Database" = db_name(DBID), "Table" = object_name(ObjectID, DBID),
      IndexID, RowsInserted, RowsDeleted, RowsUpdated, LockWaits
from monOpenObjectActivity
order by RowsInserted desc
```

- **What makes a table “hot”?**
 - Logical reads?
 - Physical reads?
 - Number of queries?
 - Lock usage?
- **monOpenObjectActivity reports a number of measures of table and index activity**
- **Example**

```
select * into #t
from master..monOpenObjectActivity
go
```

```
select TableName = object_name(ObjectID, DBID), IndexID,
       LogicalReads, PhysicalReads, Operations, LockWaits
from #t
order by 3, 4
go
```

TableName	IndexID	LogReads	PhysReads	Operations	LockWaits
-----	-----	-----	-----	-----	-----
products_tb	0	282294	9043	609	97
products_tb	2	36450	0	0	0
cust_tab	0	12315	0	17	2
cust_tab	2	239	0	0	0

Understanding and Using Historical Tables

Using Historical Tables

- **Which MDA tables are “historical” tables?**
- **What are Historical Tables?**
- **How do they work?**
- **What is the correct size to configure them?**
- **Archiving historical table data**
- **Tips on using historical tables**

Which Tables are Historical Tables?

- **monSysSQLText**
 - Records every SQL command executed on the server
- **monSysPlanText**
 - Records the Query Plan for every SQL command executed on the server
- **monSysStatement**
 - Reports the statistics for every statement within every query, batch, stored procedure, trigger, etc. executed on the server
- **monErrorLog**
 - Records every row written to the server errorlog
- **monDeadLock**
 - Records information on every deadlock that occurs on the server

What are Historical Tables?

- **The historical MDA tables contain a record of “events” within the ASE**
 - E.g., SQL submitted for a query, a statement executed within a batch, error message added to the errorlog
- **The data for these tables is stored in memory in fixed-sized arrays**
 - Size is configurable using sp_configure
- **Data in Historical tables is transient**
 - The arrays are managed as “ring buffers”: After the last entry in the array is written the first entry will be overwritten
- **Historical tables are “stateful.”**
 - The ASE remembers which records a process has already seen
 - Subsequent queries on same table will return only new records
- **Why are they stateful?**
 - This allows applications to accurately collect or “drain” the rows in these tables without finding duplicates.

- The ASE maintains the connection's *currency* in the MDA table
- Currency is reset for each new connection

8:22 AM: rows 1 – 4
Are inserted...

“Select * from monSysSQLText”

8:27 AM
User 2:
Never queried table
before. Will see
all rows.

8:26 AM: rows 5 - 8
Are inserted...

Row 1
Row 2
Row 3
Row 4
Row 5
Row 6
Row 7
Row 8

8:24 AM
User 1:
First query returns rows
1 - 4.

8:30 AM
User 1:
Second query will
return only rows added
since last query (5 – 8).

- **These sp_configure parameters determine the number of rows in the historical tables**
 - errorlog pipe max messages
 - plan text pipe max messages
 - sql text pipe max messages
 - statement pipe max messages
 - deadlock pipe max messages
- **The value of the parameter is the number of rows *per engine***
- **Correct size depends on**
 - Rate at which rows are written to table
 - Frequency with which queries will be run against the table
- **For example:**
 - 2 engines
 - 5000 rows per minute per engine
 - Select * from monSysStatement every 5 minutes
 - Statement pipe max messages should be greater than or equal to 25000
 - Result set size??? (50000 rows!)
- **Errorlog and deadlock pipes are usually much smaller than plan text, sql text and statement pipes**

Rate x Frequency = Size
E.g.: 5000/min x 5 min = 25000

Reasonable size on
busy system??
Could be >> 100000

Tips on Using Historical Tables

- **Do not use in subqueries or joins**
- **Save contents of tables to an archive table or database for analysis**
- **When collecting long-term data, archive data on a regular basis and size tables to avoid data loss**
- **How do you know whether the table for the buffer has wrapped?**
 - If # of rows returned = size of buffer * # of engines
 - In other words, if you get the entire size of the buffer, some rows were probably lost
 - Only a “rule of thumb”
 - Currently, it is not possible to determine how many rows were lost

Archiving Historical Table Data

Archiving Historical Table Data

- **Why is an historical archive useful?**
- **Because data in historical tables is transient**
 - Capture data for later or detailed analysis
- **Because repeated queries on historical tables will not return the same rows**
 - Even in subqueries or joins
 - Makes analytical queries directly against MDA tables difficult
- **Solution**
- **Data from historical tables should be moved to separate storage for analysis**
- **Create a monitoring data repository for historical diagnostics or capacity planning**

- **A possible approach: a 'collector' stored proc which frequently extracts data from the MDA tables**
 - 'sp_mda_collect'
 - sp_mda_collect 'start' [, 'hh:mm:ss'] -- runs in a loop (default interval = **30 sec.**)
 - sp_mda_collect 'stop' -- run from a different session, stops the original procedure
 - sp_mda_collect 'status' -- displays #rows saved in archive tables
 - (download from www.sypron.nl/mda)
 - Uses a separate database to collect the historical data in permanent tables
 - The permanent tables have the same layout as the historical MDA tables
 - Added a composite unique index **with ignore_dup_key** on key columns (SPID, KPID, etc.) to filter out duplicates (in case the proc needs to be restarted...)

Enterprise Monitoring Repository and Center

- **To access MDA tables from a remote server**
 - Create the MDA proxy tables on a central server
 - Map MDA proxy tables to each monitored server
- **Reduces load on monitored ASE servers**
- **Provides central source of monitoring data for your enterprise**
- **Allows easy archiving of enterprise data to permanent storage in database on repository server**

Creating an Enterprise Monitoring Center

- **Create monitoring database on central server**
- **Copy and edit installmontables script**
 - Two options:
 - Create separate monitoring database for each monitored server
 - Add server name to MDA table names to create unique table names for each server within a single database
- **Set the use database command to use the correct database**
- **Change the “loopback” server name to the remote server name of the monitored server in your central server**

Modifying installmontables Script

- **Creating MDA proxy tables in a separate database for each monitored server**

```
...  
...  
use monitor_svrtest1  
go  
...  
...  
  
create existing table  
monProcedureCache (  
    Requests          int,  
    Loads             int,  
    Writes            int,  
    Stalls            int,  
)  
external procedure  
at "svrtest1...$monProcedureCache"  
go
```

Use a separate database for each Monitored server

Proxy table points to monitored server

- **Creating MDA proxy tables in a single database for all monitored server**

```
...  
...  
use monitordb  
go  
...  
...
```

Database in which all proxy tables
Will be created

```
create existing table  
monProcedureCache_svrtest1 (  
    Requests int,  
    Loads int,  
    Writes int,  
    Stalls int,  
)  
external procedure  
at "svrtest1...$monProcedureCache"  
go
```

Unique table name constructed by a
Appending server name

Proxy table points to monitored
server

Performance Impact of MDA Tables

Performance Impact of MDA Tables

- **Two questions**
 - Impact of data collection?
 - Impact of querying MDA tables?
- **General performance impact: 5% or less**
- **Depends on a number of factors**
 - Configuration of server (e.g., number of engines, memory size, processor speed)
 - Load on server
 - Configuration of Monitoring parameters
- **Different monitoring configuration settings have different performance impacts**
- **Fully enabling all options will have greatest impact**

Performance Impact: Configuration Settings

- **Lowest impact**
 - Enable monitoring with no other options
- **Tables enabled**
 - monEngine, monDataCache, monProcedureCache, monOpenDatabases, monSysWorkerThread, monNetworkIO, monLocks, monCachePool, monIOQueue, monDeviceIO, monProcessWorkerThread, monProcessNetIO
- **Low Impact Parameters**
 - wait event timing
 - plan text pipe active
 - sql text pipe active
 - errorlog pipe active
 - deadlock pipe active
- **Greatest impact**
 - Per Object Statistics
 - Statement Historical table and Process Statements
- **Tables enabled**
 - monOpenObjectActivity, monProcessObject, monProcessActivity, monSysStatement, monProcessStatement
- **High Impact Parameters**
 - statement pipe active
 - statement statistics active
 - per object statistics
 - statement pipe active

Understanding and Handling Counter Wrap

What is Counter Wrap?

- **All MDA counter columns are 32-bit signed integers**
 - Maximum value is 2147483647
- **When signed integers are incremented above maximum value they become negative**
 - $2147483647 + 1 \Rightarrow -2147483646$
- **Internal adjustments prevent MDA counter values from becoming negative**
 - Therefore counter ranges are from 0 to 2147483647
- **When the ASE increments an MDA counter past the maximum value it will return to 0 and start increasing again**

Handling Counter Wrap with Delta Values

- If counter has wrapped, add difference between start value and maximum value + 1 to the current value of the counter

```

Select CacheName,
       CacheSearches =
       case
         when e.CacheSearches < s.CacheSeaches
         then
           (2147483648 - s.CacheSearches) + e.CacheSearches)
         else
           (e.CacheSearches - s.CacheSearches)
       end
from #cacheStart s, #cacheEnd e
where s.CacheID = e.CacheID
    
```

- Again: As long as change in counter values is ≤ 2147483647 , delta values will be accurate

Which MDA Table Columns Can Wrap?

- **Not all MDA columns are likely to wrap**
 - Some counter values increment slowly
 - Some numeric columns are not counters
- **Columns that can wrap pretty quickly**
 - **monDataCache**
 - **CacheSearches**
 - **LogicalReads**
 - **monNetworkIO**
 - **BytesSent**
 - **BytesReceived**
 - **monSysWaits**
 - **Waits**
- **Others wrap less quickly**
 - **monEngine.ContextSwitches**
 - **monNetworkIO.PacketsSent**

sp_sysmon, subqueries, joins...

Using MDA Tables and sp_sysmon

- **Monitor Counters are a set of counters used by sp_sysmon and Monitor Server**
- **Some MDA table columns are derived from Monitor Counters**
- **sp_sysmon resets the value of Monitor Counters when it starts**
- **This can have an impact on applications using MDA tables or Monitor Server**
- **MDA table columns that come from Monitor Counters are documented.**
 - **Attributes column = “counter, reset”**

Using MDA Tables and sp_sysmon

- **The ASE 12.5.3 release introduced changes to sp_sysmon that allow it to run without clearing monitor counters**
- **Also enhanced so that when multiple applications are using monitor counters the collection of monitor data will not be terminated until all applications are finished**
- **It is safe to run sp_sysmon when using the MDA tables as long as sp_sysmon is run with the 'noclear' option.**

Subqueries, Joins and Self-Joins

- **Rule of thumb: Don't use joins or subqueries when querying the MDA tables**
- **Why? Because the MDA table data is transient and reflects the ASE's instantaneous state, joins and subqueries may not give the expected result.**
 - Sequential queries on same table can give different results
- **Because of the currency mechanism, self-joins or subqueries involving one of the historical tables more than once will not work.**
 - Currency is reset by first query and the same rows will not be seen by the subquery or inner join table
- **Solution: Copy MDA table data to a work table or permanent repository when complex analysis is required.**

Analyzing Stored Procedure Performance

Analyzing Stored Procedure Performance

- **Historical Server** provides stored procedure performance information
- **MDA tables** do not provide a table with historical stored procedure statistics
- The **monSysStatement** table can be used to report this information

```
select ProcName =
    isnull(object_name(ProcedureID, DBID),
    "UNKNOWN"),
    DBName = isnull(db_name(DBID), "UNKNOWN"),
    ElapsedTime = datediff(ms, min(StartTime),
    max(EndTime))
from master..monSysStatement
group by SPID, DBID, ProcedureID, BatchID
having ProcedureID != 0
```

```

1> select ProcName = isnull(object_name(ProcedureID, DBID), "UNKNOWN"),
2> DBName = isnull(db_name(DBID), "UNKNOWN"),
3> ElapsedTime = datediff(ms, min(StartTime), max(EndTime))
4> from master..monSysStatement
5> group by SPID, DBID, ProcedureID, BatchID
6> having ProcedureID != 0
7> order by 3
8> go

```

ProcName	DBName	ElapsedTime
-----	-----	-----
p_sybugstatus	engcomdb	1096
sybrev_fetch_revstatus	engcomdb	983
p_sybugstatus	engcomdb	923
p_sybugstatus	engcomdb	836
p_sybugstatus	engcomdb	683
p_sybugstatus	engcomdb	620
p_sybugstatus	engcomdb	586
p_sybugstatus	engcomdb	543
p_sybugstatus	engcomdb	533
p_sybugstatus	engcomdb	526
...		
...		
0		

- **Aggregate performance statistics can be derived from the output of the previous query**

```
/*
** Build a detail table
*/
select ProcName = isnull(object_name(ProcedureID, DBID),
    "UNKNOWN"),
    DBName = isnull(db_name(DBID), "UNKNOWN"),
    ElapsedTime = datediff(ms, min(StartTime), max(EndTime))
into #t1
from master..monSysStatement
group by SPID, DBID, ProcedureID, BatchID
having ProcedureID != 0

/*
** Calculate aggregate values
*/
select ProcName, DBName, "Avg" = avg(ElapsedTime),
    NumExecs = count(*)
from #t1
group by ProcName, DBName
order by 3 desc
```


- Determine average elapsed time and total executions for stored procedures

```
1> select ProcName, DBName, "AvgElapsed" = avg(ElapsedTime),
2> NumExecs = count(*)
3> from #t1
4> group by ProcName, DBName
5> order by 3 desc
6> go
```

ProcName	Database	AvgElapsed	NumExecs
-----	-----	-----	-----
p_sybbugstatus	engcomdb	483	32
sn_temp_filters_qts1	qts_db	330	26
sy_resolution_insert	qts_db	260	44
p_sybbugreleasematrix	engcomdb	186	21
create_sn_subscriptions	qts_db	108	9
p_sybbugsrelease	engcomdb	91	37
sn_temp_filters_qts2	qts_db	83	2
sn_temp_filters_qts4	qts_db	73	11
sn_get_next_key	qts_db	69	5
create_sn_filters	qts_db	65	5

...

...

- Identify statements within stored procedures consuming greater than average elapsed time

```
/*
** Build work table
*/
select ProcName = isnull(object_name(ProcedureID, DBID), "UNKNOWN"),
       DBName = convert(char(15), isnull(db_name(DBID), "UNKNOWN")),
       LineNumber,
       ElapsedTime = datediff(ms, StartTime, EndTime)
into #t1
from master..monSysStatement
where ProcedureID != 0
/*
** Calculate aggregate values and find problematic statements
*/
select ProcName, DBName, LineNumber, "AvgElapsed" = avg(ElapsedTime)
from #t1
group by DBName, ProcName, LineNumber
having avg(ElapsedTime) > (select avg(ElapsedTime) from #t1)
order by 4 desc
```

...

...

ProcName	DBName	LineNumber	AvgElapsed
-----	-----	-----	-----
row_update	qts_db	614	2160
p_sybbugstatus	engcomdb	60	240
row_update	qts_db	147	98
row_insert	qts_db	308	98
e2_CiMember	qts_db	71	77
p_sybbugstatus	engcomdb	56	76
sy_addl_case_update	qts_db	138	70
p_sybbugstatus	engcomdb	125	69
sybrev_report_newcrs	engcomdb	48	30
log_activity	qts_db	155	18
p_sybbugstatus	engcomdb	29	16
p_sybbugstatus	engcomdb	145	15
log_activity	qts_db	90	14

...

...

```

1> select * into #t1 from master..monSysStatement
2> go
1> select ProcedureName = isnull(object_name(ProcedureID, DBID), "UNKNOWN"),
2> "Database" = db_name(DBID),
3> "Execs" = count(*)
4> from #t1
5> where ProcedureID != 0
6> group by DBID, ProcedureID
7> order by 3 desc
8> go

```

ProcedureName	Database	Execs
-----	-----	-----
sp_mltypeset	empdb	8138
p_sybugstatus	engcomdb	888
sp_help_rep_agent	sybsystemprocs	462
p_sybugsrelease	engcomdb	205
sn_get_next_key	qts_db	176
create_sn_filter_criteria	qts_db	162
create_sn_subscriptions	qts_db	136
create_sn_filters	qts_db	120
...		
...		

Recent Enhancements

Enhancements in 12.5.1

- **360 columns in 12.5.0.3 (first version of MDA tables)**
- **5 new columns in 12.5.1**
 - **monErrorLog.State** - state of error
 - **monOpenDatabases.QuiesceTag** - tag specified with quiesce database (if any)
 - **monOpenDatabases.SuspendedProcesses** - number of currently suspended processes due to log-full condition in this database
 - **monProcessWorkerThread.FamilyID** - spid of parent process
 - **monProcessWorkerThread.ParallelQueries** - total # parallel queries attempted

- **2 new columns in 12.5.2**
 - **monProcessObject.TableSize** - table size in Kb
 - **monProcessActivity.WorkTables** - total number of work tables created by the process

- **Fixes:**
 - milliseconds fixed in **monSysStatement.StartTime / EndTime**
 - can be used to determine the exact duration of each statement (resolution = 3 milliseconds)

Enhancements in 12.5.3

- **4 new columns in 12.5.3**
 - **monProcessActivity.ServerUserID** - Login ID (suid)
 - **monProcessSQLText.ServerUserID** - Login ID (suid)
 - **monSysSQLText.ServerUserID** - Login ID (suid)
 - **monProcessProcedures.LineNumber** – Line number in a stored procedure
- **Fixes:**
 - Various small bugs were fixed



Your Questions are Welcome

Thanks!

robv@sybase.com

peter.dorfman@sybase.com