

Getting Started with SAP Sybase IQ Column Store Analytics Server

Lesson 7: Row-Level Versioning



Table of Contents

1. Introduction.....	1
2. How RLV Works.....	2
2.1 Architecture.....	2
2.2 Memory Storage.....	2
2.3 Lock Management.....	3
3. RLV Tutorial.....	4
3.1 Configuration	4
3.2 Create Table and Load Data.....	5
3.3 Merge	10
3.4 Monitor Write-Intent Locks	11
4. Summary.....	17

1. Introduction

In your SAP Sybase IQ database, load, edit or delete data performance is sometimes sacrificed for column storage and increased analytical speed. The row level versioning (RLV) data store can be used to process row-level updates, inserts and deletes in real time, which is useful for improving the performance in some cases.

RLV data storage, introduced in SAP Sybase IQ 16, is optimized for writes. A table can be RLV enabled, which allows multiple transactions to modify different rows of same table concurrently. This minimizes lock contention and allows for low latency writes.

In this lesson, you will learn how the row level versioning data store works, as well as how to monitor locks. This is a useful store for writing data to tables, but the required memory usage needs to be considered.



Speed up the server!

2. How RLV Works

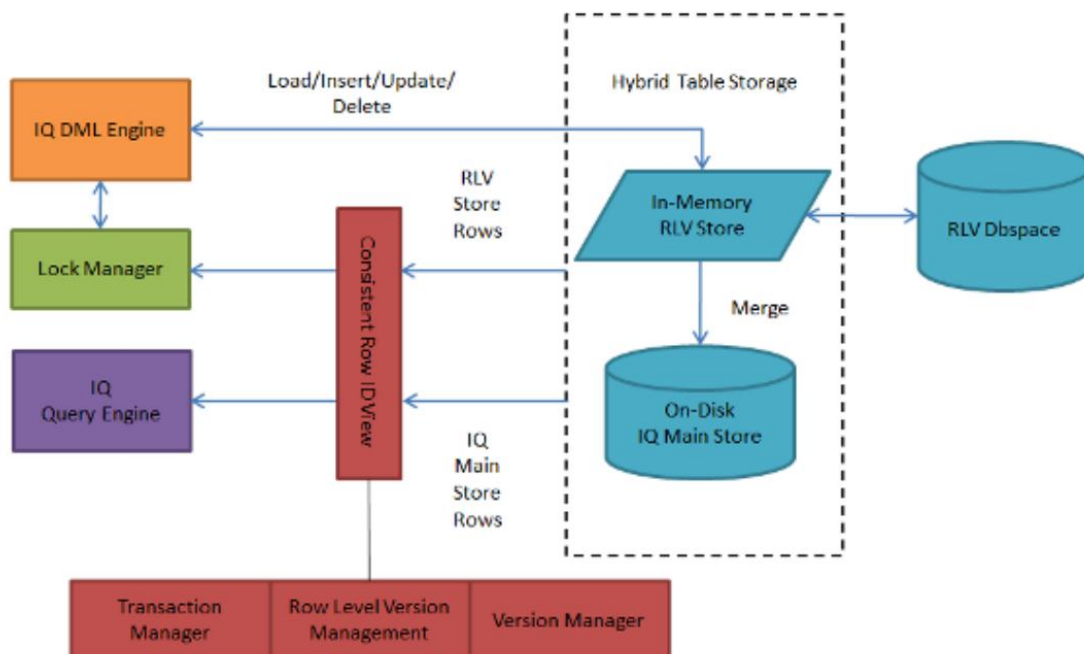
In-memory row-level versioning is a hybrid storage model which optimizes writing data. For any RLV enabled table, when data is loaded or written it is stored in the RLV store in real-time, which has a preconfigured size. The RLV store is later merged into the IQ main store when it reaches a threshold or at a time limit. Then, memory is freed from the RLV store, and the newly written data can be used in the IQ main store for analytics.

Note: RLV can only be used in simplex configurations, at the current time.

2.1 Architecture

Row-level storage is enabled in a hybrid table storage, which is made up of RLV and IQ main stores. Logically, the data is stored in the single hybrid table storage, but the server will track the actual data location when querying and manipulating data.

The transaction manager controls the flow of transactions, blocking and allowing them as required. When a single table-level versioning (TLV) transaction is writing, the transaction manager will block all other writes to the table. During row-level versioning transactions, however, only other writes to that row are blocked.



2.2 Memory Storage

When starting the database, memory is set aside for the RLV store. This memory is completely independent of the IQ main cache. Any RLV enabled table is able to write rows to that store, until it is full.

The size of the RLV store should be determined based on the number of tables, load frequency and merge frequency. For most systems, 250-500MB per table marked for each RLV enabled table is

sufficient. This can be set by the “-iqlvmem” startup flag. If this flag is not specified, the value defaults to 2GB.

2.3 Lock Management

In SAP Sybase IQ, locks occur automatically. They are used to ensure consistency between concurrent transactions. Without RLV, as in versions of IQ earlier than 16 or tables without RLV enabled, when a table is locked no other users can have write access to the entire table, but they can gain read access. All of a single transaction’s locks are held by the server until a commit or rollback.

In RLV enabled tables, locks are handled differently to allow concurrent writes to the same table, on different rows. These RLV enabled tables have schema locks, row locks and write-intent locks. These locks are defined as follows:

- Schema locks lock the schema of the table, preventing transactions from being disturbed by a difference in columns or other schema changes.
- Write-intent locks are held on tables, if they intend to write to some row in the future. Many requestors can hold a write-intent lock at the same time, as this lock does not imply current writes. A prerequisite to this lock is that the table can be accessed by read-write transactions with row-level snapshot versioning. All write-intent locks must be released before locks on the entire table can be held.
- Row locks allow the holder to write to any column of the row. These locks can only be obtained if the table already has a write-intent lock. This lock is used for updating or deleting rows, as there can only be one holder of the lock on each row at a time.

3. RLV Tutorial

In this tutorial, we will configure our environment to allow for row-level versioning. We can then create a table which is flagged for row-level versioning, load data and see the memory usage. We will then look at locks on the database as transactions are executed on RLV-enabled tables.

3.1 Configuration

Open “tpch.cfg”, in the /opt/sybase/TPCHDB directory (or similar directory, as chosen in previous lessons). Add the following line:

```
-iqlvmem 500
```

This gives the RLV store 500MB of memory. Now, when we run our server, it can store up to 500MB in the RLV store.

Start the IQ server, by running the following command in the terminal:

```
start_iq @tpch.cfg tpch.db
```

Now open Sybase Control Center:

```
$SYBASE/SCC-3_2/bin/scc.sh
```

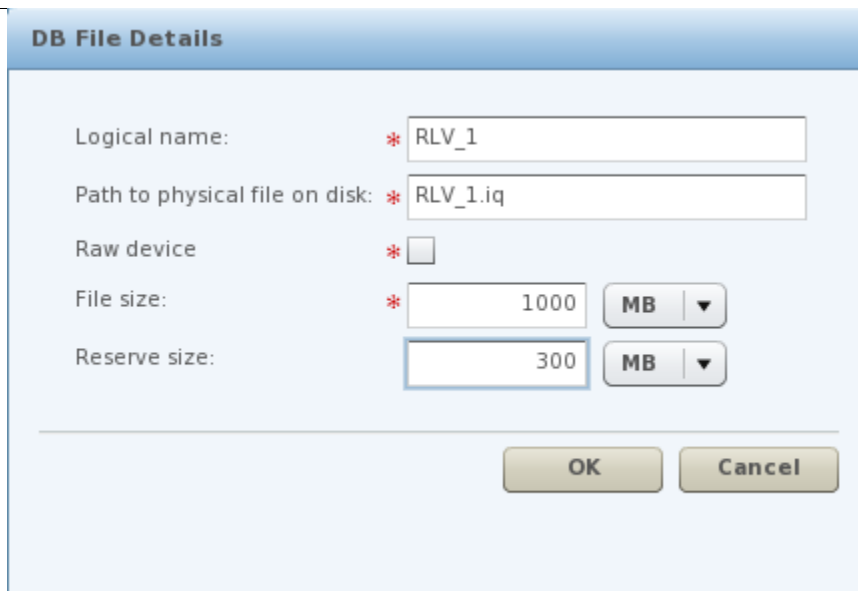
From SCC, we will create a RLV dbspace, in a similar way to the dbspace created in Lesson 3, Create Schema and Load Data.

Open the Administration Console for the “tpch” server. Then, expand IQ Servers, followed by Space Management. Click the drop-down for Dbspaces and click “New...”.

Add a new dbspace in RLV store.

The screenshot shows the 'Create Dbspace Wizard' window. On the left, a tree view shows 'DB Files' under 'General Details'. The main area is titled 'General Details' and contains the following fields: 'Resource' (dropdown menu set to 'tpch'), 'Name of dbspace to create' (text field with 'RLV'), 'Store' (radio buttons for 'Main', 'Catalog', and 'RLV', with 'RLV' selected), 'Striping' (radio buttons for 'On' and 'Off', with 'On' selected), and 'Stripe size(Kb)' (text field with '1'). At the bottom right are four buttons: 'Back', 'Next', 'Finish', and 'Cancel'.

Click “Next”. Create a dbfile, with size 1000MB.

A screenshot of a 'DB File Details' dialog box. It has a light blue header with the title 'DB File Details'. Below the header, there are five fields, each with a red asterisk indicating it is required. The first field is 'Logical name:' with the value 'RLV_1'. The second field is 'Path to physical file on disk:' with the value 'RLV_1.iq'. The third field is 'Raw device' with an unchecked checkbox. The fourth field is 'File size:' with the value '1000' and a unit dropdown menu set to 'MB'. The fifth field is 'Reserve size:' with the value '300' and a unit dropdown menu set to 'MB'. At the bottom right of the dialog are 'OK' and 'Cancel' buttons.

DB File Details

Logical name: * RLV_1

Path to physical file on disk: * RLV_1.iq

Raw device * ☐

File size: * 1000 MB ▼

Reserve size: 300 MB ▼

OK Cancel

Click “OK” and then “Finish” on the resulting screen.

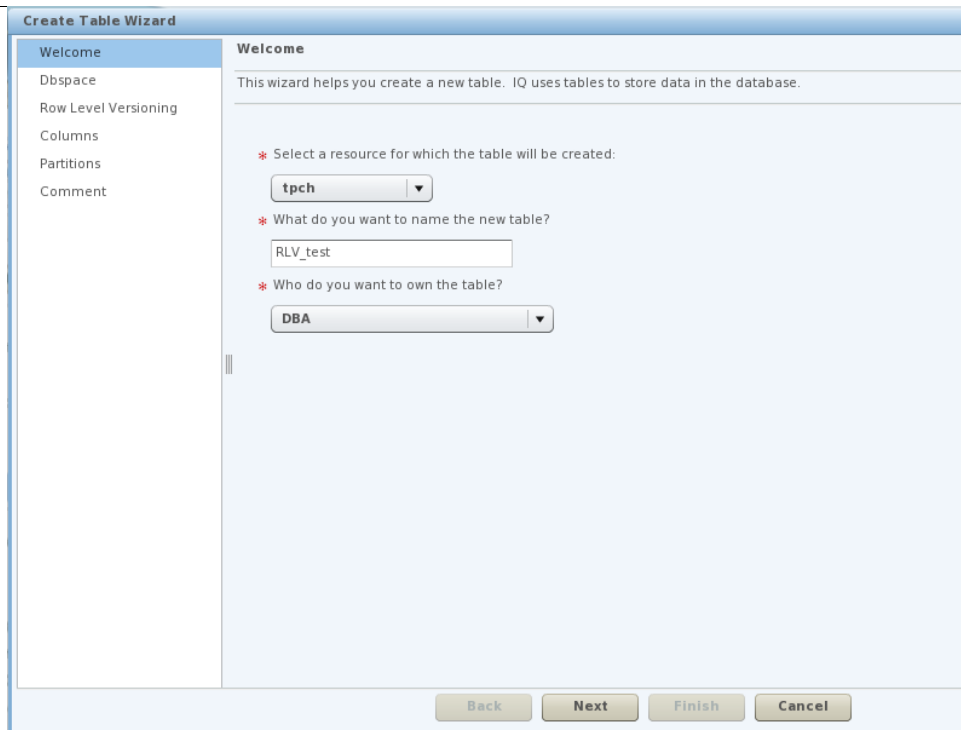
The RLV dbspace, as just created, holds the persistent logs and data from the tables which have RLV storage enabled. The reserve size is used by data structures during table operations, such as load.

3.2 *Create Table and Load Data*

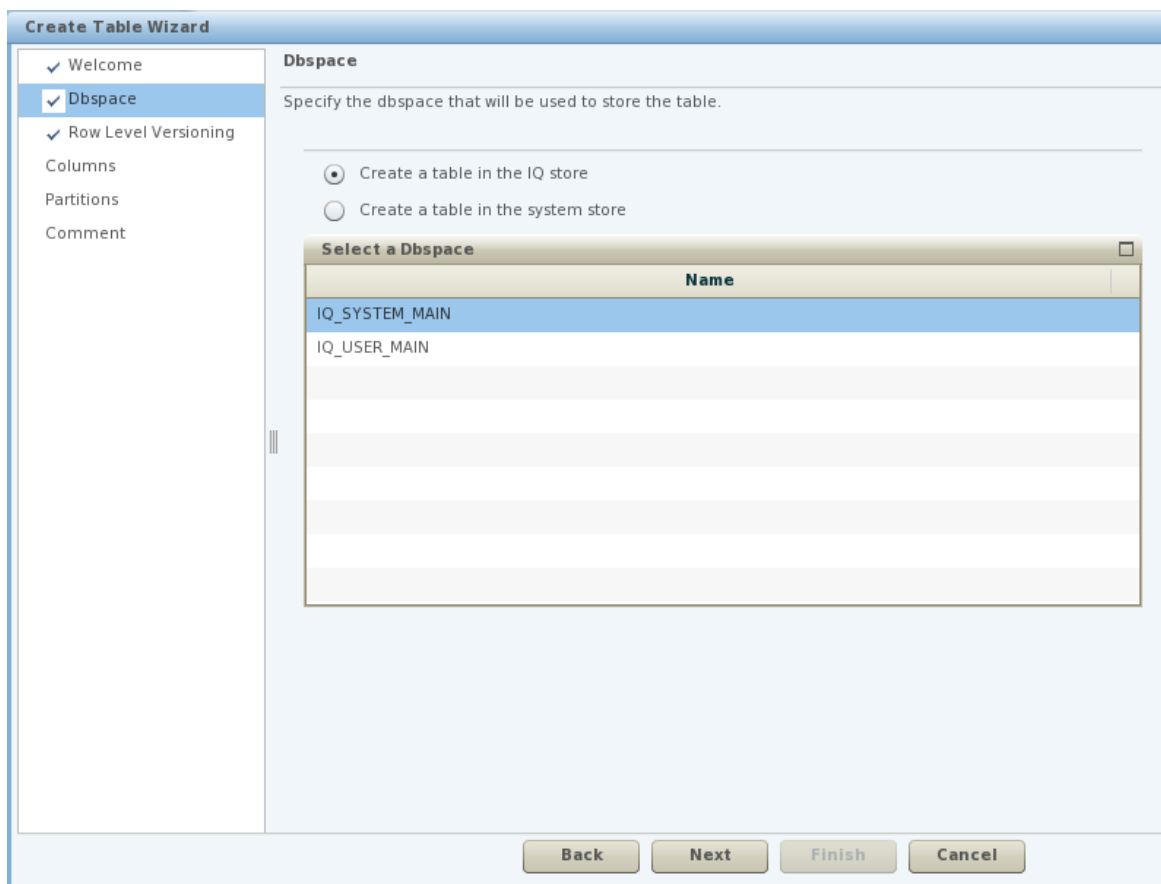
Any given table in the database must be flagged for RLV operation for the write-optimized storage to be used.

Note: RLV cannot be enabled on a table with foreign keys. Catalog, temporary and global temporary tables are also not supported.

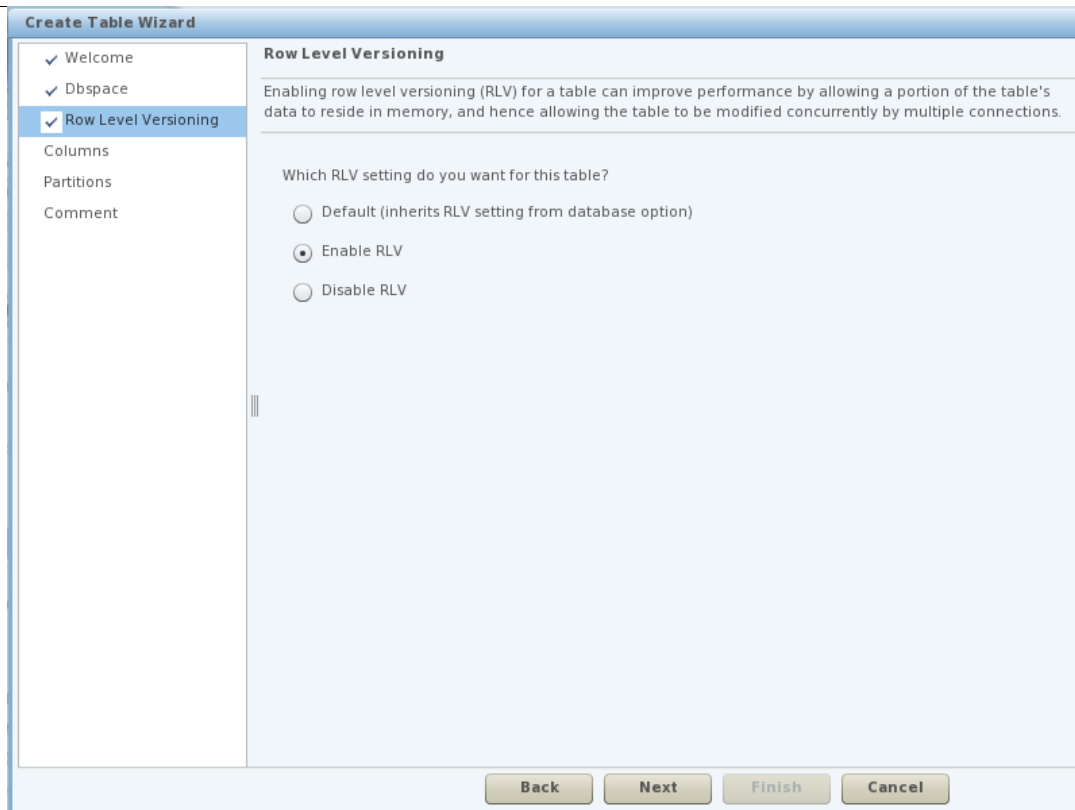
In the Administration Console, expand IQ Servers, followed by Schema Objects and Tables. Then click the drop-down beside “Tables” and click “New...”. Follow the wizard by filling out the following fields:



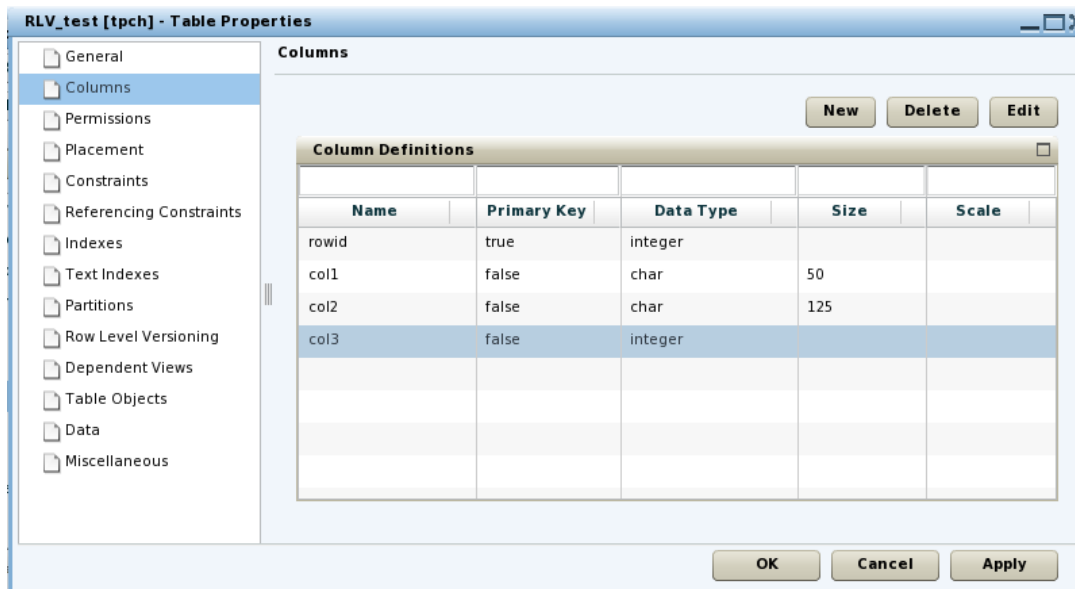
Click “Next”. Create the table in the “IQ_SYSTEM_MAIN” dbspace.



Click “Next”. Enable RLV.



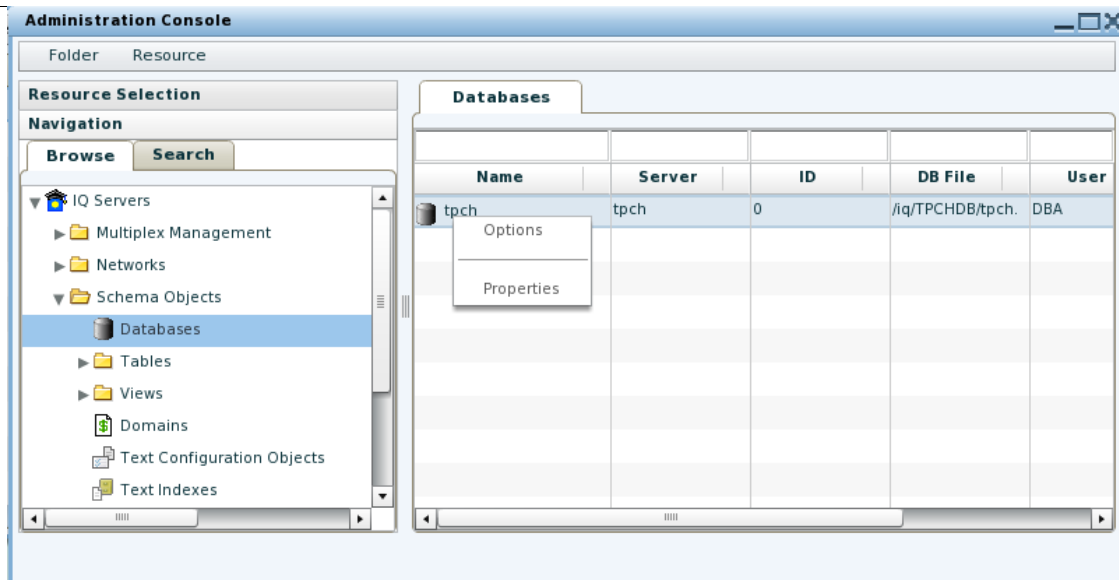
Click “Next”. Use the column wizard to add some test columns, all stored in “IQ_SYSTEM_MAIN”



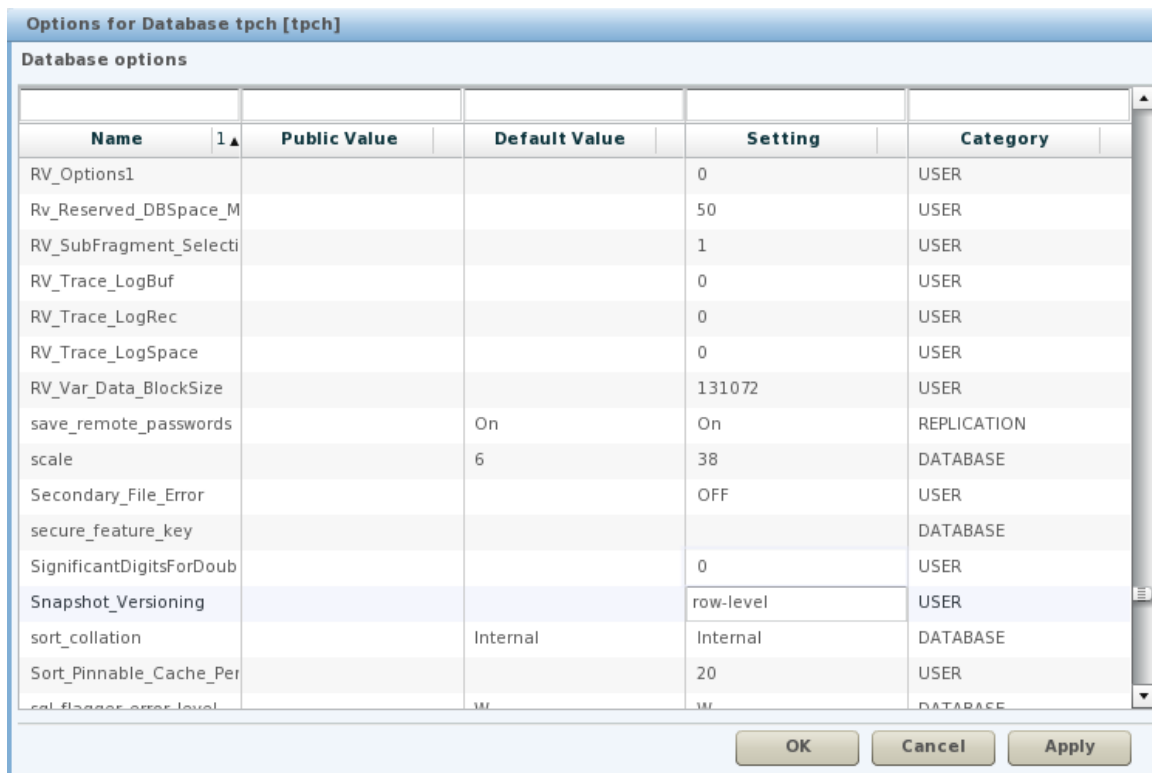
Click “Next”, “Next” and “Finish”, not adding any partitions.

Now, we will set snapshot versioning to row-level.

In the left pane of the Administration Console, expand IQ Servers, Schema Objects and click on Database. In the drop-down beside the database name (tpch, in this example), click “Options”.



In the options window, scroll down to find “Snapshot_Versioning” and change the value to “row-level”



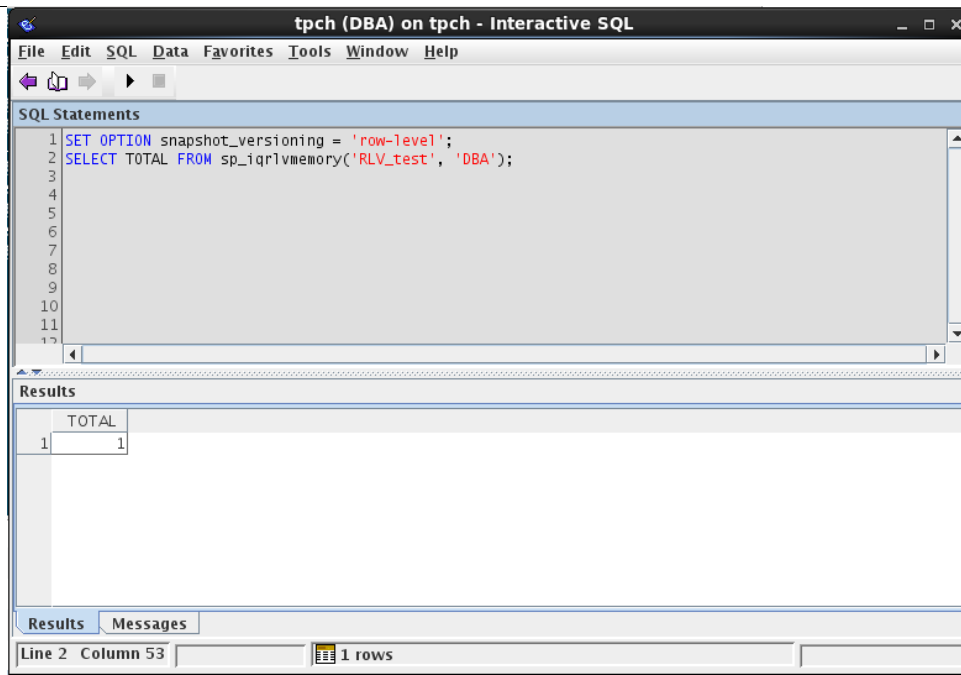
Click “Apply” and “OK”

We will perform the next steps in Interactive SQL. In a terminal window, execute the following command:

```
dbisql -datasource "tpchdb"
```

Check memory usage before data is loaded into the RLV table.

```
SELECT TOTAL FROM sp_iqrlvmemory('RLV_test', 'DBA');
```

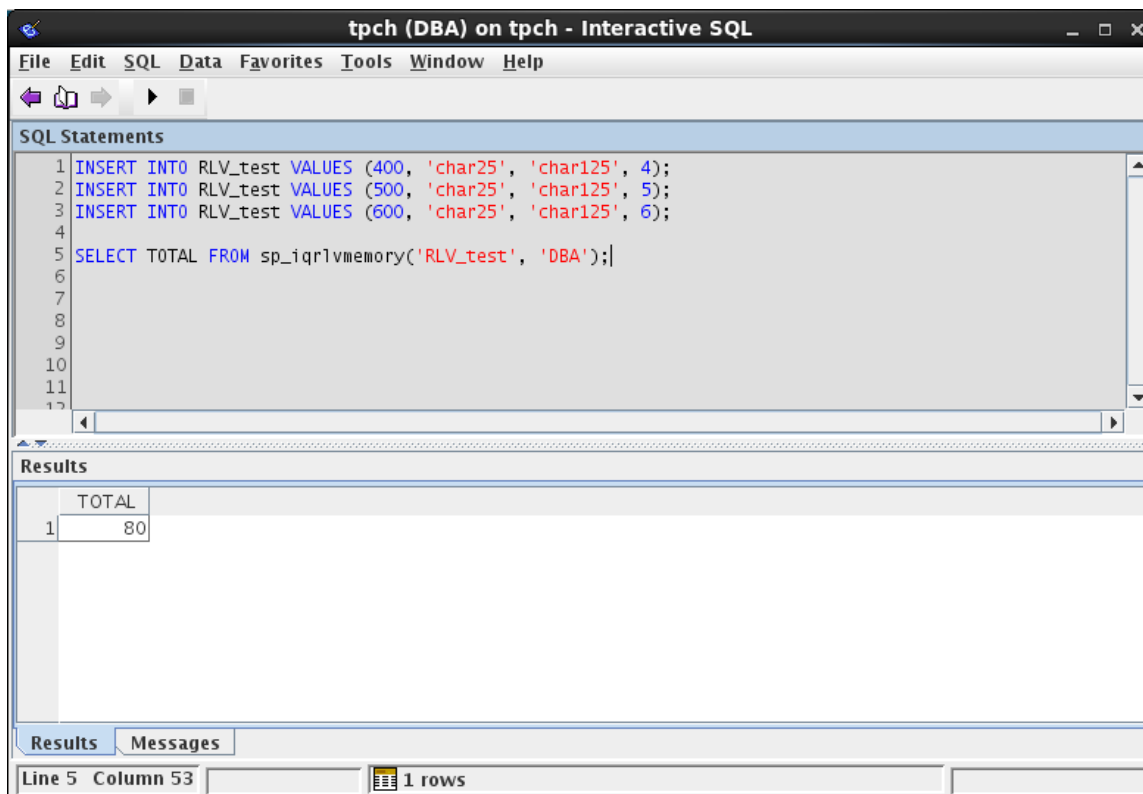


Insert data into your created table, changing values based on created columns.

```
INSERT INTO RLV_test VALUES (1, 'char25', 'char125', 4);
INSERT INTO RLV_test VALUES (2, 'char25', 'char125', 5);
INSERT INTO RLV_test VALUES (3, 'char25', 'char125', 6);
```

Check memory usage again, after the data has been loaded.

```
SELECT TOTAL FROM sp_iqrlvmemory('RLV_test', 'DBA');
```



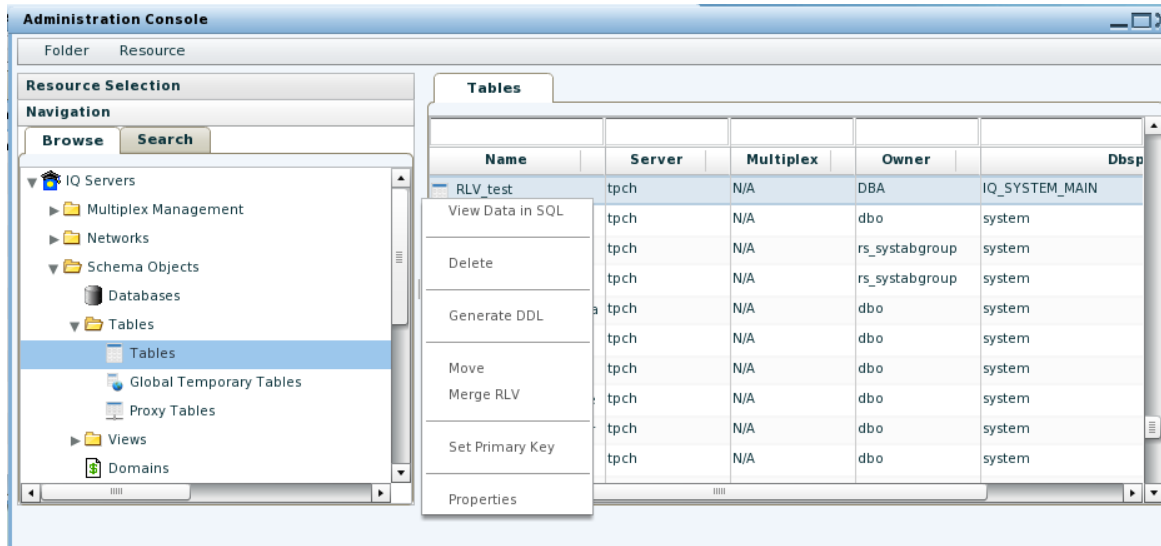
Notice that the total memory in the RLV memory store has increased, as the inserted rows were put in the RLV memory store.

3.3 Merge

The RLV memory will be automatically merged into IQ main store based on time or memory threshold. You can also manually force a merge.

Close Interactive SQL. We do this, to stop any transactions on the RLV_test table.

If not still open, open Sybase Control Center and go to the Tables section of the Administration Console. Then, click the drop-down to the right of the RLV table, “RLV_test”, in this example. Choose “Merge RLV”.



In the pop-up, choose “Yes” to proceed.

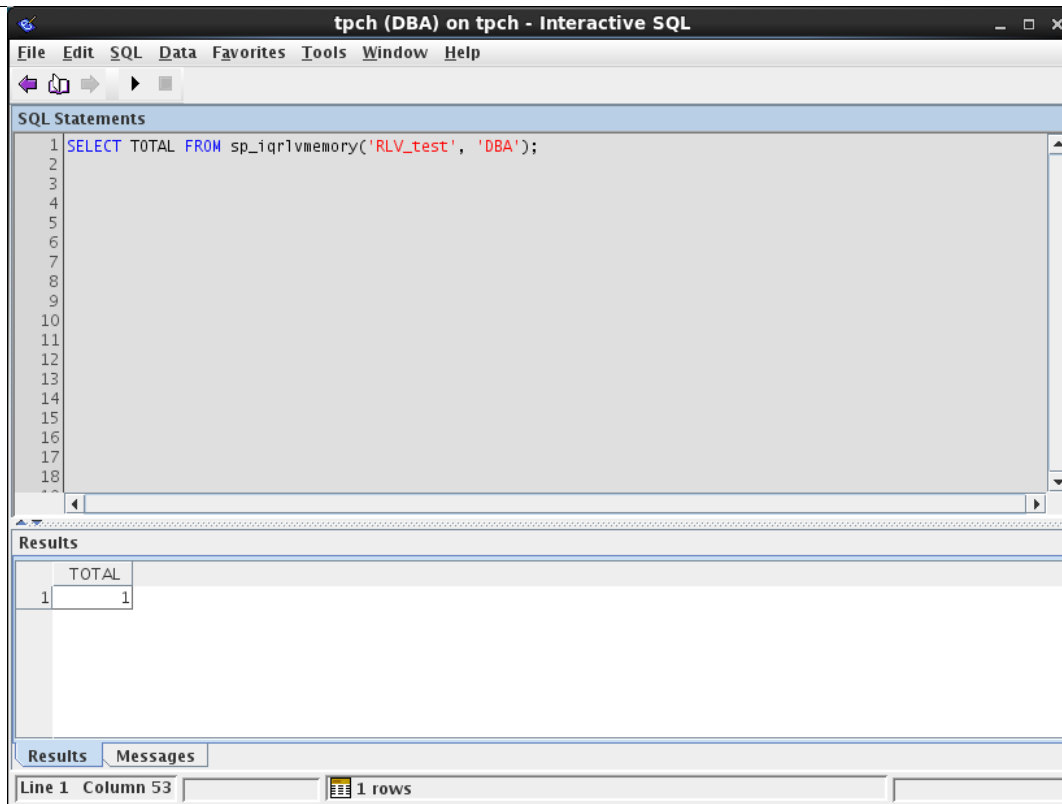
We will now reopen Interactive SQL to check the RLV store memory again.

Execute the following command in a terminal window:

```
dbisql -datasource "tpchdb"
```

Run the following command, to check RLV memory usage again:

```
SELECT TOTAL FROM sp_iqrlvmemory('RLV_test', 'DBA');
```

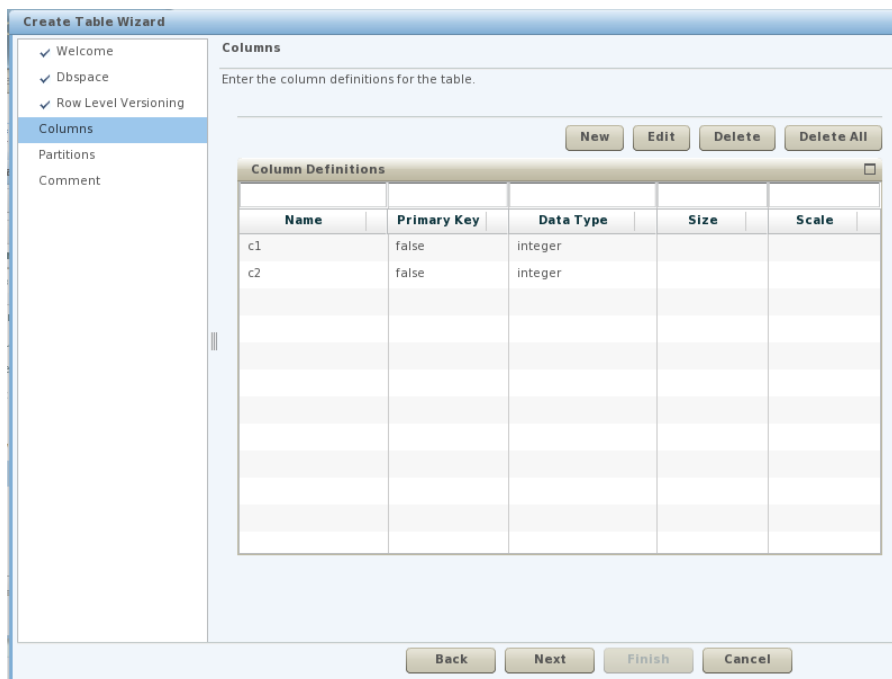


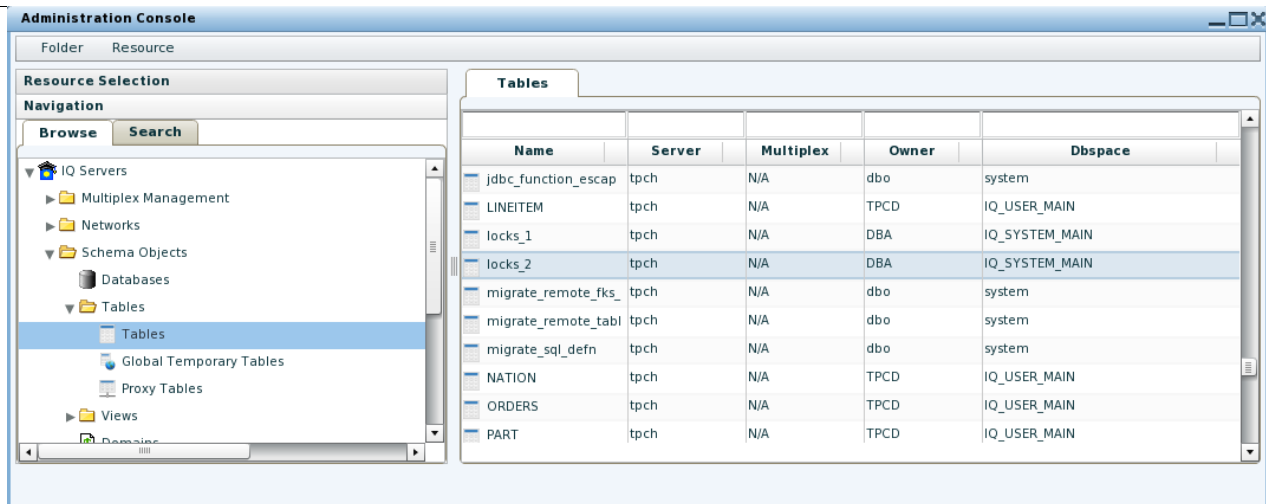
Notice the total memory has gone back down.

3.4 *Monitor Write-Intent Locks*

Change “Snapshot_Versioning” to “table_level” in the Administration Console of SCC, as done before, by changing the options of the database.

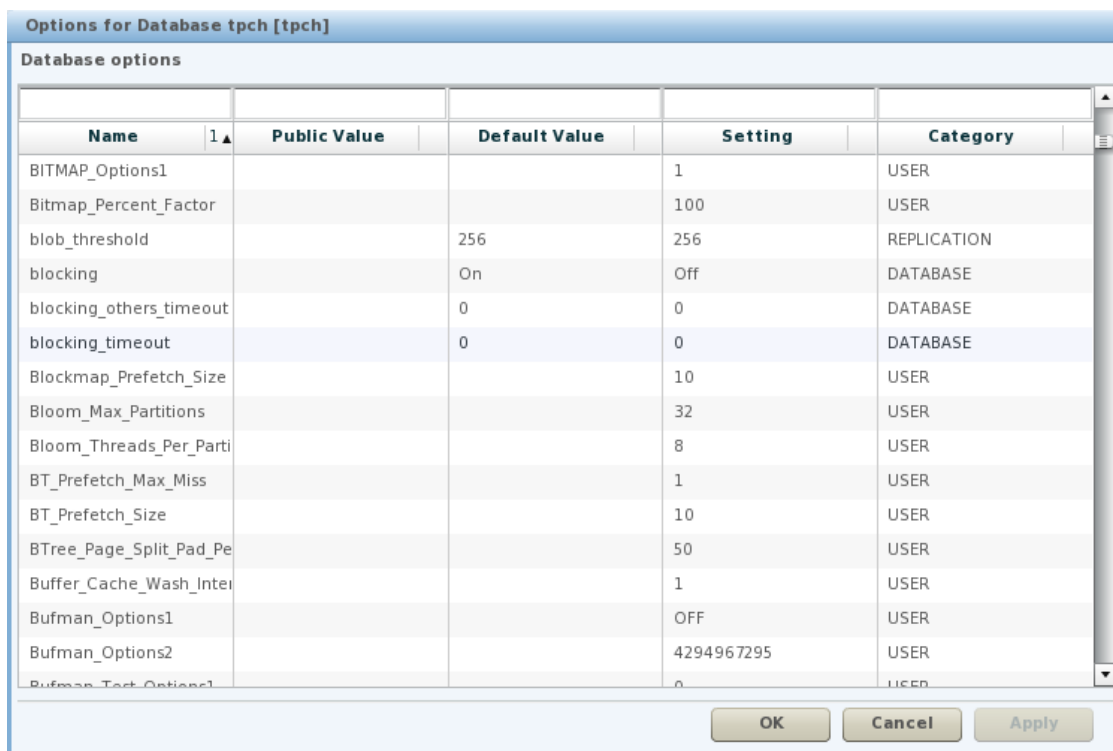
Create two tables, as earlier, each with RLV enabled. Create columns as you wish.





Check the database properties as before, ensuring that “blocking” is on and “blocking_timeout” is 0. Blocking controls the behavior in response to locking conflicts.

When blocking is enabled, blocks occur once a transaction has obtained a lock. This means that once a transaction obtains a lock any other transaction that attempts to obtain that lock must wait, until the original transaction has released the lock or the blocking timeout has been reached. When blocking is disabled, errors occur if a transaction attempts to obtain an already obtained lock.



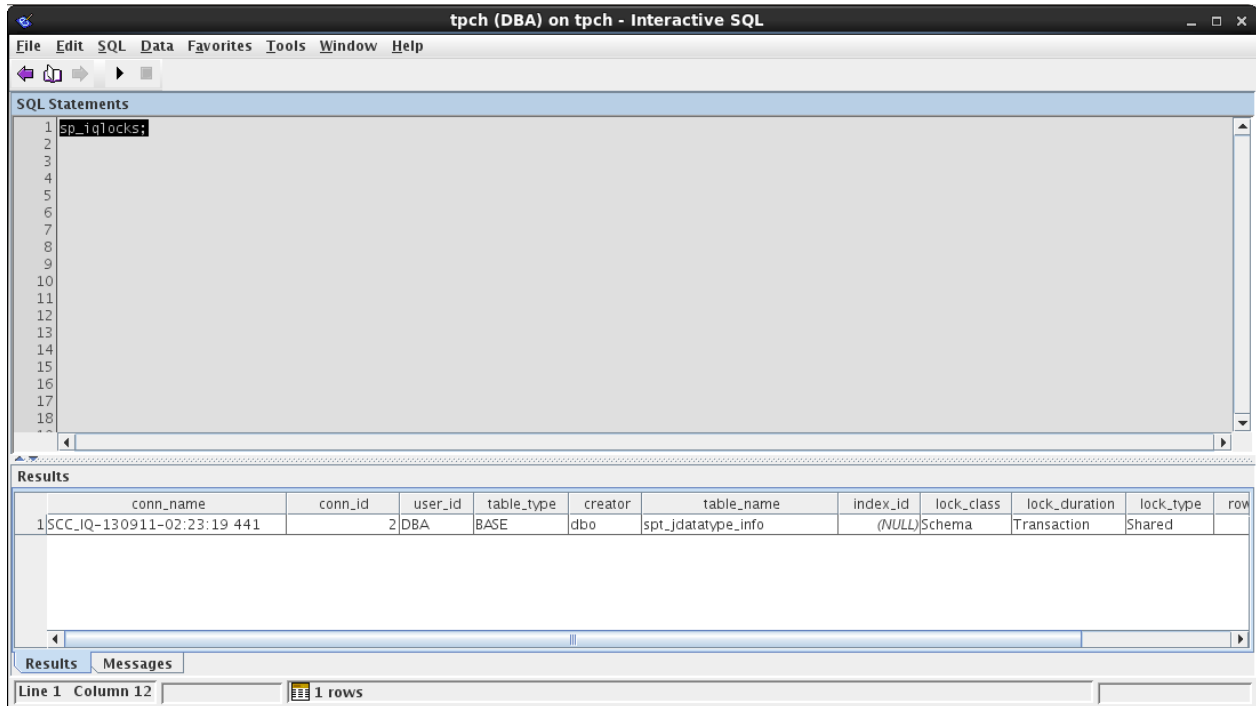
We will now view the current set of database locks, ensuring that no write-intent locks are obtained yet (and thus, are not displayed).

Now, open Interactive SQL, as earlier, running the following command in a terminal window.

```
dbisql -datasource "tpchdb"
```

Run the following SQL statement and click “Execute”.

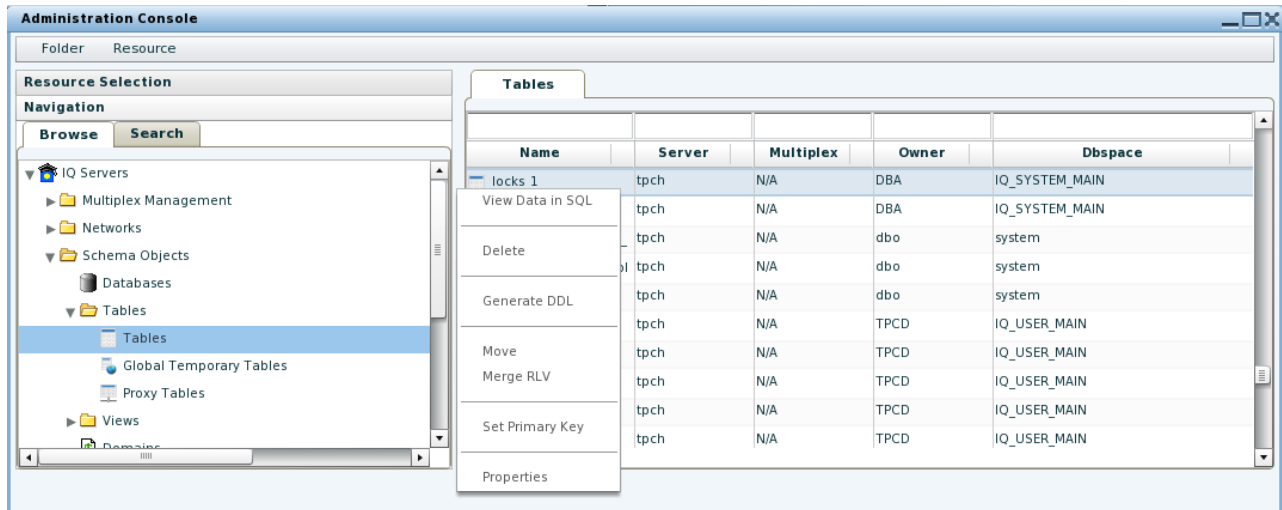
sp_iqlocks;



Notice that there are no write-intent locks or locks on the newly created lock_1 or lock_2 table.

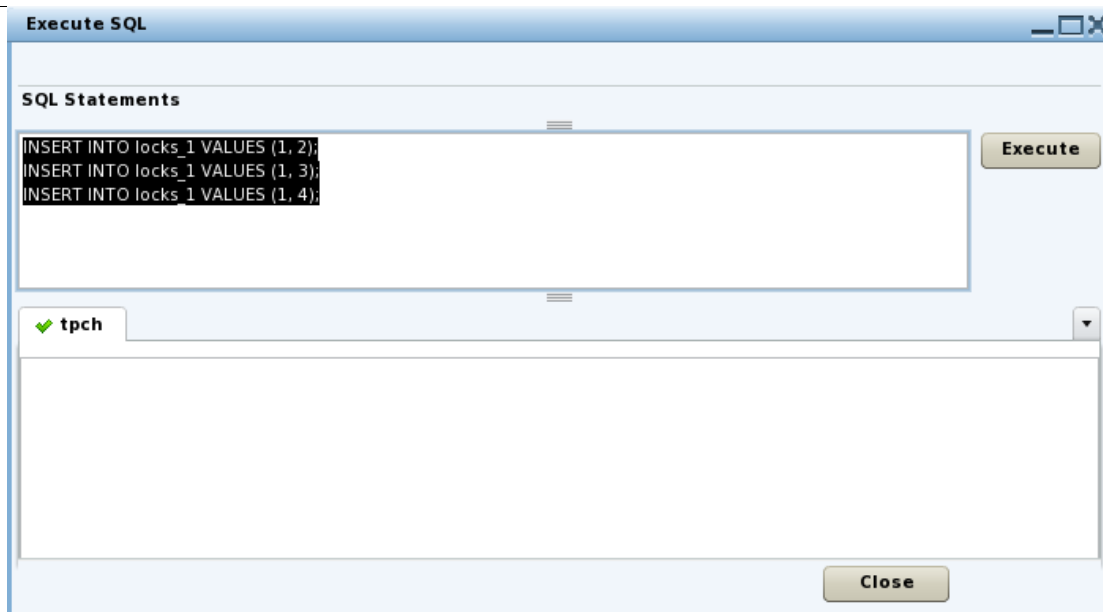
A write intent lock will have “lock_type” set to “Write”. Each lock has a specified “table_name”, which tells you what table the lock is on.

Now, close this window and go back to the Administration Console. View the options for the database, as done in “3.2 Create Table and Load Data” above. Change the “Snapshot_Versioning” value to “row-level”. Click the drop-down arrow and choose “View Data in SQL”



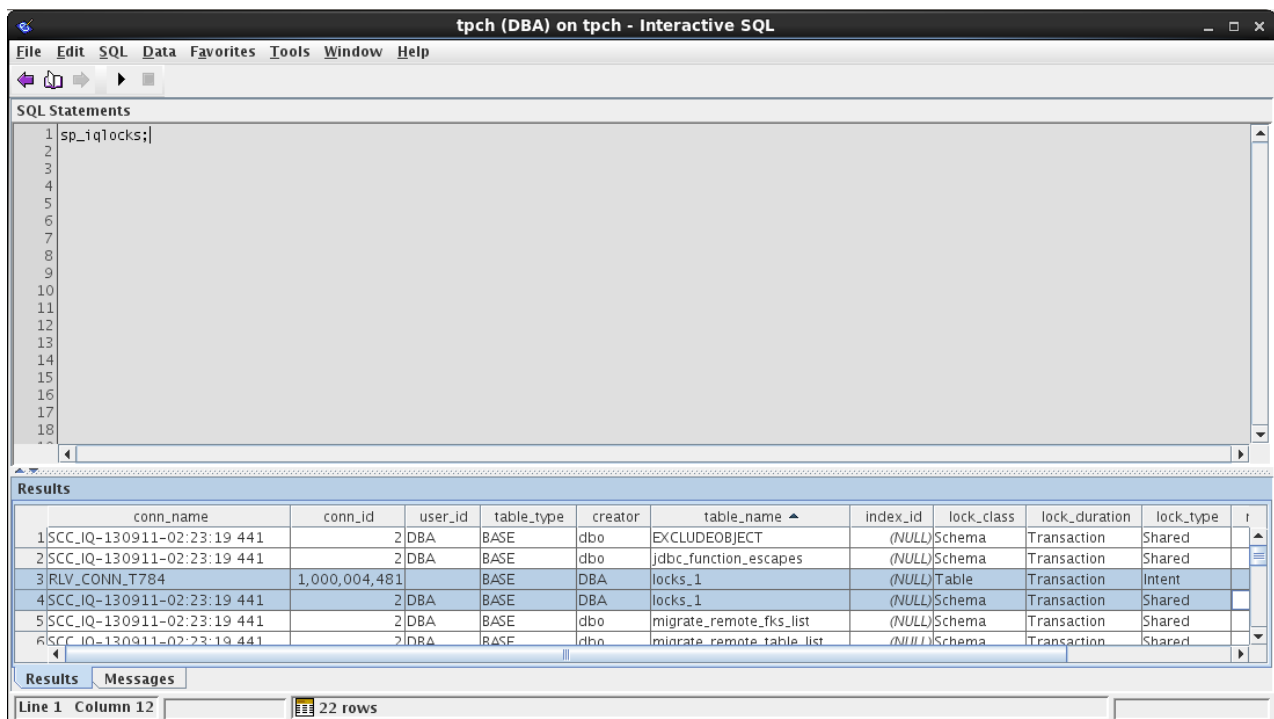
Change the values in the SQL Statements textbox to the following, and execute.

```
INSERT INTO locks_1 VALUES (1, 2);  
INSERT INTO locks_1 VALUES (2, 2);  
INSERT INTO locks_1 VALUES (3, 2);
```



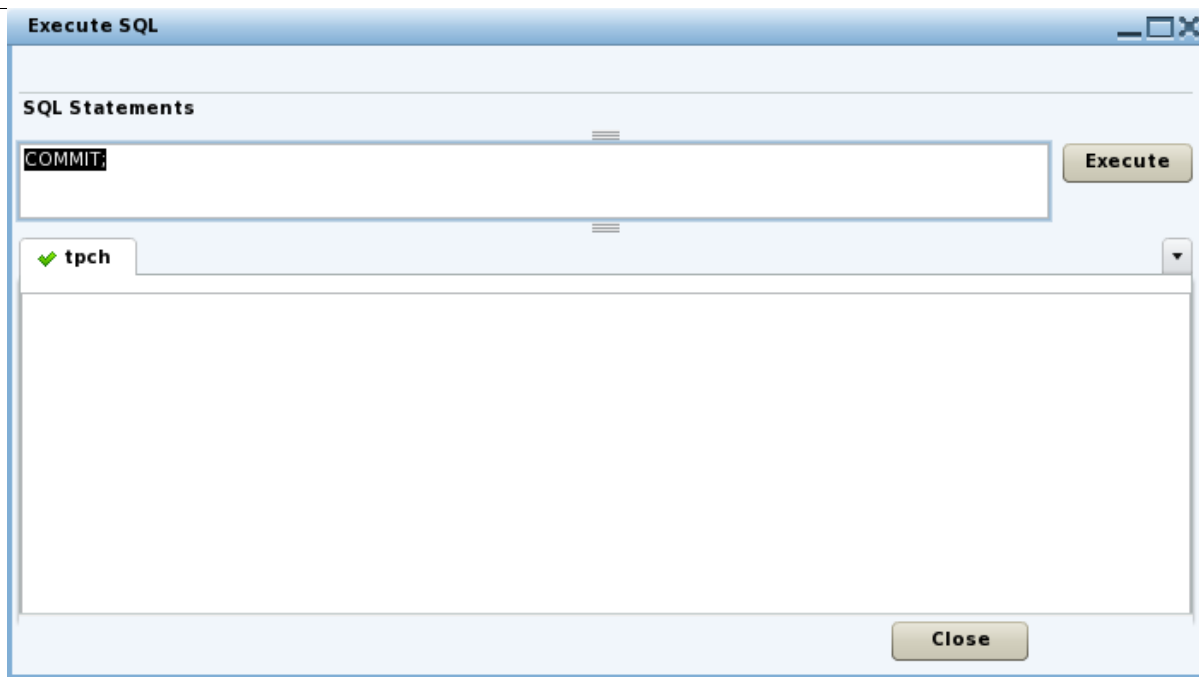
This creates the RLV-enabled portion of the table in memory.

Run “sp_iqlocks” in Interactive SQL again.



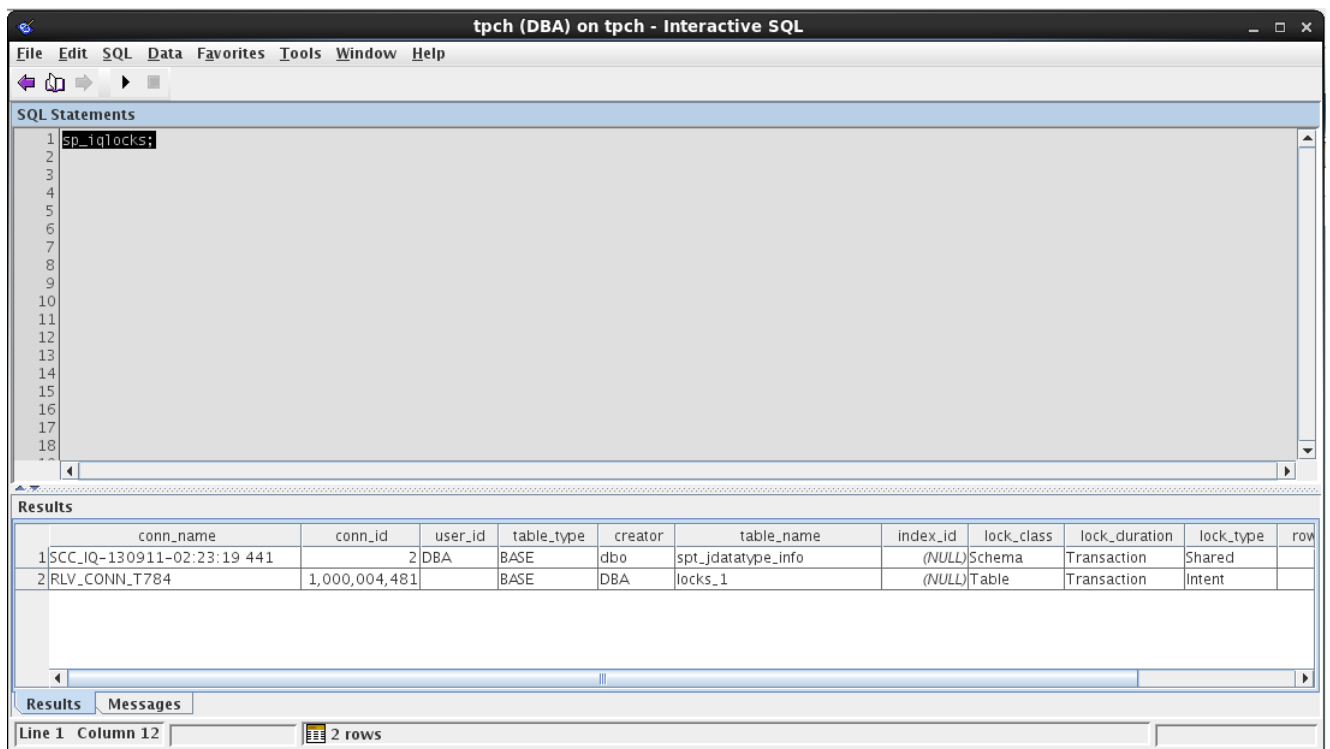
You will see a schema lock and a table lock for the locks_1 table, with the uncommitted insert statements.

Go back to the Administration Console in SCC, and in the “View Data in SQL” pop-up of the “locks_1” table, run “COMMIT;” Commit releases the locks held by the transaction, in this case, it will be the schema lock.



Now, run “sp_iqlocks” in Interactive SQL again.

You will now see that the schema lock is gone. The rows inserted in the table are now in the RLV-enabled portion of the table, and this write-intent lock is held by that part of the table.



These locks can also be viewed, in less detail, in the “Connections” tab of the “Monitor Node” of the “tpch” server in SCC.

Overview

All Statistics

Engine

Connections

Transactions

Dbspaces

Store I/O

Caches

Table Versions

Operations & Requests

Network

Table Pages

Table Locks

Deadlocks

Settings

Connections

Connection Statistics

Active Connections

User ID	Name	Creation time	Connection ID	State	Client IP Ad...	Connection ...
DBA	SCC_IQ-130911-02:	09/11/2013 2:23 AM	21	Active		CONNECTION
DBA	SQL_DBC_1f83170	09/11/2013 4:42 AM	4211	Active		CONNECTION
	INT: RLVLockConn	09/11/2013 4:51 AM	4483	Active	NA	CONNECTION

Connection Details

Associated Transactions

Statement being executed:

NO COMMAND

Last request time:

Request type:

unknown (0)

Last command time:

12/30/0 12:00 AM

Command type:

NONE

Transaction ID	Creation time	State

Data load completed.

Last refresh: Wed Sep 11 2013 04:55:20 AM

4. Summary

In this chapter you have learned about the new row-level versioning and how it works with SAP Sybase IQ. You have done a lesson, in which you created a table which performs write operations in the RLV memory store and merges the data to IQ main store. You can now use in-memory row-level versioning to optimize writes, when applicable.

More information and tutorials can be found at:

<http://infocenter.sybase.com/help/topic/com.sybase.infocenter.dc01840.1601/doc/pdf/iqrlv.pdf>

© 2013 SAP AG. All rights reserved.

SAP, R/3, SAP NetWeaver, Duet, PartnerEdge, ByDesign, SAP BusinessObjects Explorer, StreamWork, SAP HANA, and other SAP products and services mentioned herein as well as their respective logos are trademarks or registered trademarks of SAP AG in Germany and other countries.

Business Objects and the Business Objects logo, BusinessObjects, Crystal Reports, Crystal Decisions, Web Intelligence, Xcelsius, and other Business Objects products and services mentioned herein as well as their respective logos are trademarks or registered trademarks of Business Objects Software Ltd. Business Objects is an SAP company.

Sybase and Adaptive Server, iAnywhere, Sybase 365, SQL Anywhere, and other Sybase products and services mentioned herein as well as their respective logos are trademarks or registered trademarks of Sybase Inc. Sybase is an SAP company.

Crossgate, m@gic EDDY, B2B 360°, and B2B 360° Services are registered trademarks of Crossgate AG in Germany and other countries. Crossgate is an SAP company.

All other product and service names mentioned are the trademarks of their respective companies. Data contained in this document serves informational purposes only. National product specifications may vary.

These materials are subject to change without notice. These materials are provided by SAP AG and its affiliated companies ("SAP Group") for informational purposes only, without representation or warranty of any kind, and SAP Group shall not be liable for errors or omissions with respect to the materials. The only warranties for SAP Group products and services are those that are set forth in the express warranty statements accompanying such products and services, if any. Nothing herein should be construed as constituting an additional warranty.

